

A modern conference room with large windows and a long table. The room is empty, with several office chairs arranged around the table. The view outside the windows shows a cityscape. The image has a warm, orange-toned overlay.

MySQL Aggregate Functions

A modern office interior with large windows and a long conference table. The room is brightly lit, and the windows offer a view of a city skyline. The conference table is long and dark, with several chairs arranged around it. The text "COUNT()" is overlaid in the center of the image.

COUNT()

COUNT()

So far:

COUNT()

So far:

- Theory of Relational Databases
- SQL Theory
- Coding Techniques and Best Practices
- SELECT, INSERT, UPDATE, DELETE

COUNT()

So far:

- Theory of Relational Databases
- SQL Theory
- Coding Techniques and Best Practices
- SELECT, INSERT, UPDATE, DELETE

Next:

COUNT()

So far:

- Theory of Relational Databases
- SQL Theory
- Coding Techniques and Best Practices
- SELECT, INSERT, UPDATE, DELETE

Next:

- Aggregate Functions

COUNT()

aggregate functions

they gather data from *many* rows of a table, then aggregate it into a *single* value

COUNT()

aggregate functions

they gather data from *many* rows of a table, then aggregate it into a *single* value

INPUT

COUNT()

aggregate functions

they gather data from *many* rows of a table, then aggregate it into a *single* value

INPUT

the information contained
in *multiple* rows

COUNT()

aggregate functions

they gather data from *many* rows of a table, then aggregate it into a *single* value

INPUT

the information contained in *multiple* rows →

COUNT()

aggregate functions

they gather data from *many* rows of a table, then aggregate it into a *single* value

INPUT

the information contained
in *multiple* rows



OUTPUT

COUNT()

aggregate functions

they gather data from *many* rows of a table, then aggregate it into a *single* value

INPUT

the information contained
in *multiple* rows



OUTPUT

the *single* value they
provide

COUNT()

COUNT()

COUNT()



COUNT()

SUM()

COUNT()



COUNT()

SUM()

MIN()

COUNT()



COUNT()

SUM()

MIN()

MAX()

COUNT()



COUNT()

SUM()

MIN()

MAX()

AVG()

COUNT()

COUNT()

SUM()

MIN()

MAX()

AVG()

COUNT()

COUNT()

SUM()

MIN()

MAX()

AVG()

aggregate functions

COUNT()

COUNT()

SUM()

MIN()

MAX()

AVG()

aggregate functions

=

summarizing functions

COUNT()

- Why do these functions exist?

COUNT()

- Why do these functions exist?

- they are a response to the information requirements of a company's different organizational levels

COUNT()

- Why do these functions exist?

- they are a response to the information requirements of a company's different organizational levels
- top management executives are typically interested in *summarized* figures and rarely in detailed data

COUNT()

COUNT()

applicable to both *numeric* and *non-numeric* data

COUNT()

- **COUNT(DISTINCT)**

helps us find the number of times unique values are encountered in a given column

COUNT()

- aggregate functions *typically* ignore null values throughout the field to which they are applied

COUNT()

- aggregate functions typically ignore null values throughout the field to which they are applied

COUNT()

- aggregate functions typically ignore null values throughout the field to which they are applied



COUNT()

- aggregate functions typically ignore null values throughout the field to which they are applied



only if you have indicated a *specific* column name within the parentheses

COUNT()

- aggregate functions typically ignore null values throughout the field to which they are applied



only if you have indicated a *specific* column name within the parentheses

Alternatively:

COUNT()

- aggregate functions typically ignore null values throughout the field to which they are applied



only if you have indicated a *specific* column name within the parentheses

Alternatively:

- COUNT(*)

COUNT()

- aggregate functions typically ignore null values throughout the field to which they are applied



only if you have indicated a *specific* column name within the parentheses

Alternatively:

- `COUNT(*)`

* returns the number of all rows of the table, `NULL` values *included*

COUNT()

COUNT()

COUNT()

COUNT()

the parentheses and the argument must be attached to the name of the aggregate function

COUNT()

COUNT()

the parentheses and the argument must be attached to the name of the aggregate function

- you shouldn't leave white space before opening the parentheses

COUNT()

COUNT()

the parentheses and the argument must be attached to the name of the aggregate function

- you shouldn't leave white space before opening the parentheses

COUNT()

COUNT()

COUNT()

the parentheses and the argument must be attached to the name of the aggregate function

- you shouldn't leave white space before opening the parentheses

COUNT()

COUNT ()

COUNT()

COUNT()

the parentheses and the argument must be attached to the name of the aggregate function

- you shouldn't leave white space before opening the parentheses

COUNT()

COUNT_()

COUNT()

COUNT()

the parentheses and the argument must be attached to the name of the aggregate function

- you shouldn't leave white space before opening the parentheses

COUNT()

~~COUNT_()~~

A modern office interior with large windows and a long conference table. The room is brightly lit, and the view outside the windows shows a cityscape. The text "SUM()" is centered in the image.

SUM()

SUM()

COUNT(*)

SUM()

COUNT(*)

* returns all rows of the table, NULL values *included*

SUM()

COUNT(*)

* returns all rows of the table, NULL values *included*

SUM(*)

SUM()

● COUNT(*)

* returns all rows of the table, NULL values *included*

● SUM(*)



SUM()

COUNT(*)

* returns all rows of the table, NULL values *included*

~~SUM(*)~~

* goes well with only the COUNT() function

SUM()

COUNT()

SUM()



COUNT()

- applicable to both *numeric* and *non-numeric* data

SUM()

COUNT() - applicable to both *numeric* and *non-numeric* data

SUM()

MIN()

MAX()

AVG()

SUM()

COUNT() - applicable to both *numeric* and *non-numeric* data

SUM()

MIN()

MAX()

AVG()

SUM()

COUNT() - applicable to both *numeric and non-numeric* data

SUM()

MIN()

MAX()

AVG()

- work *only with numeric* data

MIN() and MAX()

MIN() and MAX()

MAX()

returns the maximum value of a column

MIN() and MAX()

- **MAX()**

returns the maximum value of a column

- **MIN()**

returns the minimum value of a column

A modern office interior with large windows and a long conference table. The room is brightly lit, and the windows offer a view of a city skyline. The conference table is long and dark, with several chairs arranged around it. A laptop is open on the table. The text "AVG()" is centered in the image.

AVG()

AVG()

AVG()

extracts the average value of all non-null values in a field

AVG()

COUNT()

SUM()

MIN()

MAX()

AVG()

AVG()

COUNT()

SUM()

MIN()

MAX()

AVG()

- aggregate functions can be applied to any *group* of data values within a certain column

AVG()

COUNT()

SUM()

MIN()

MAX()

AVG()

- aggregate functions can be applied to any group of data values within a certain column

AVG()

COUNT()

SUM()

MIN()

MAX()

AVG()

- aggregate functions can be applied to any group of data values within a certain column

AVG()

COUNT()

SUM()

MIN()

MAX()

AVG()

- aggregate functions can be applied to any group of data values within a certain column

frequently used together with a GROUP BY clause

A modern office interior with large windows and a long conference table. The room is brightly lit, and the windows offer a view of a city skyline. The conference table is long and dark, with several chairs arranged around it. The text "ROUND()" is overlaid in the center of the image.

ROUND()

ROUND()

● ROUND(*#, decimal_places*)

ROUND()

- `ROUND(#,decimal_places)`

numeric, or math, function you can use

ROUND()

- `ROUND(#, decimal_places)`

numeric, or math, function you can use

- usually applied to the single values that aggregate functions return

A modern conference room with large windows and a long table. The room is empty, with several chairs arranged around the table. The view outside the windows shows a cityscape. The text "COALESCE() - Preamble" is overlaid on the image.

COALESCE() - Preamble

COALESCE() - Preamble

Here we will study something a bit more sophisticated.

IF NULL() and COALESCE() are among the advanced SQL functions in the toolkit of SQL professionals. They are used when null values are dispersed in your data table and you would like to substitute the null values with another value.

So, let's adjust the "Departments" duplicate in a way that suits the purposes of the next video, in which we will work with IF NULL() and COALESCE().

First, let's look at our table and see what we have there.

COALESCE() - Preamble



SQL

```
SELECT * FROM departments_dup;
```

Result Grid		Filter Rows:	Export:	Wrap Cell Content:
dept_no	dept_name			
d001	Marketing			
d002	Finance			
d003	Human Resources			
d004	Production			
d005	Development			
d006	Quality Management			
d007	Sales			
d008	Research			
d009	Customer Service			

Nine departments, with their department numbers and names provided. Ok!

COALESCE() - Preamble

Currently, as shown in the DDL statement of this table, the “Department name” field is with a NOT NULL constraint, which naturally means we must insert a value in each of its rows.

DDL for employees.departments_dup

```
1 CREATE TABLE `departments_dup` (  
2   `dept_no` char(4) NOT NULL,  
3   `dept_name` varchar(40) NOT NULL  
4 ) ENGINE=InnoDB DEFAULT CHARSET=utf8
```

COALESCE() - Preamble

Now, with the ALTER TABLE statement and the CHANGE COLUMN command, we will modify this constraint and allow null values to be registered in the “department name” column.



SQL

```
ALTER TABLE departments_dup
```

```
CHANGE COLUMN dept_name dept_name VARCHAR(40) NULL;
```

COALESCE() - Preamble

Right after that, we will insert into the department number column of this table a couple of data values – D-10 and D-11, the numbers of the next two potential departments in the “Departments Duplicate” table.



SQL

```
INSERT INTO departments_dup(dept_no) VALUES ('d010'), ('d011');
```

COALESCE() - Preamble

By running this SELECT query over here, you can see whether this operation was carried out successfully.



SQL

```
SELECT
    *
FROM
    departments_dup
ORDER BY dept_no ASC;
```

COALESCE() - Preamble

We have the two new department numbers listed below, and in the “Department name” column we can see two null values. The latter happened because we allowed for null values to exist in this field, “Department name”. Thus, Workbench will indicate that a value in a cell is missing by attaching a “null” label to it. Great!

Result Grid		Filter Rows:	Export:	Wrap Cell Content:
	dept_no	dept_name		
	d001	Marketing		
	d002	Finance		
	d003	Human Resources		
	d004	Production		
	d005	Development		
	d006	Quality Management		
	d007	Sales		
	d008	Research		
	d009	Customer Service		
	d010	NULL		
	d011	NULL		

COALESCE() - Preamble

The next adjustment we'll have to make is adding a third column called "Department manager". It will indicate the manager of the respective department. For now, we will leave it empty, and will add the NULL constraint. Finally, we will place it next to the "Department name" column by typing "AFTER "Department name".



SQL

```
ALTER TABLE employees.departments_dup  
ADD COLUMN dept_manager VARCHAR(255) NULL AFTER dept_name;
```

COALESCE() - Preamble

Let's check the state of the “Departments duplicate” table now.



SQL

```
SELECT
    *
FROM
    departments_dup
ORDER BY dept_no ASC;
```

COALESCE() - Preamble

Exactly as we wanted, right? The third column is completely empty and we have null values in the last two records. These are the “department name” and “manager” fields.

Result Grid	Filter Rows:	Export:	Wrap Cell Content:
dept_no	dept_name	dept_manager	
d001	Marketing	NULL	
d002	Finance	NULL	
d003	Human Resources	NULL	
d004	Production	NULL	
d005	Development	NULL	
d006	Quality Management	NULL	
d007	Sales	NULL	
d008	Research	NULL	
d009	Customer Service	NULL	
d010	NULL	NULL	
d011	NULL	NULL	

COALESCE() - Preamble

To save the “Departments duplicate” table in its current state, execute a COMMIT statement.



SQL

```
COMMIT;
```

Here we'll end the setup for the video about IF NULL() and COALESCE().

Good Luck!

IFNULL() and COALESCE()

IFNULL() and COALESCE()

- IFNULL(expression_1, expression_2)

IFNULL() and COALESCE()

- IFNULL(expression_1, expression_2)

returns the first of the two indicated values if the data value found in the table is *not null*, and returns the second value if there is a *null* value

IFNULL() and COALESCE()

IFNULL(expression_1, expression_2)

returns the first of the two indicated values if the data value found in the table is *not null*, and returns the second value if there is a *null* value

- prints the returned value in the column of the output

IFNULL() and COALESCE()

- COALESCE(expression_1, expression_2 ..., expression_N)

IFNULL() and COALESCE()

- `COALESCE(expression_1, expression_2 ..., expression_N)`

allows you to insert N arguments in the parentheses

IFNULL() and COALESCE()

- `COALESCE(expression_1, expression_2 ..., expression_N)`

allows you to insert N arguments in the parentheses

- think of COALESCE() as IFNULL() with more than two parameters

IFNULL() and COALESCE()

- `COALESCE(expression_1, expression_2 ..., expression_N)`

allows you to insert N arguments in the parentheses

- think of COALESCE() as IFNULL() with more than two parameters
- COALESCE() will always return a *single* value of the ones we have within parentheses, and this value will be *the first non-null value* of this list, reading the values from left to right

IFNULL() and COALESCE()

- `COALESCE(expression_1, expression_2 ..., expression_N)`

- if COALESCE() has two arguments, it will work precisely like IFNULL()

IFNULL() and COALESCE()

- **IFNULL()** and **COALESCE()** do not make any changes to the data set. They merely create an output where certain data values appear in place of NULL values.

IFNULL() and COALESCE()

- `COALESCE(expression_1, expression_2 ..., expression_N)`

IFNULL() and COALESCE()

- `COALESCE(expression_1, expression_2 ..., expression_N)`

- we can have a single argument in a given function

IFNULL() and COALESCE()

● COALESCE(~~expression_1, expression_2 ..., expression_N~~)

- we can have a single argument in a given function
- practitioners find this trick useful if some *hypothetical result* must be provided in a supplementary column

IFNULL() and COALESCE()

● COALESCE(~~expression_1, expression_2 ..., expression_N~~)

- we can have a single argument in a given function
- practitioners find this trick useful if some *hypothetical result* must be provided in a supplementary column
- COALESCE() can help you visualize *a prototype of the table's final version*

IFNULL() and COALESCE()

IFNULL()

works with precisely *two* arguments

IFNULL() and COALESCE()

IFNULL()

works with precisely *two* arguments

COALESCE()

can have *one, two, or more* arguments