

SQL Joins

A modern conference room with large windows and a long table. The room is empty, with several office chairs arranged around the table. The view outside the windows shows a cityscape. The image has a blue tint and a stylized, torn-paper-like border.

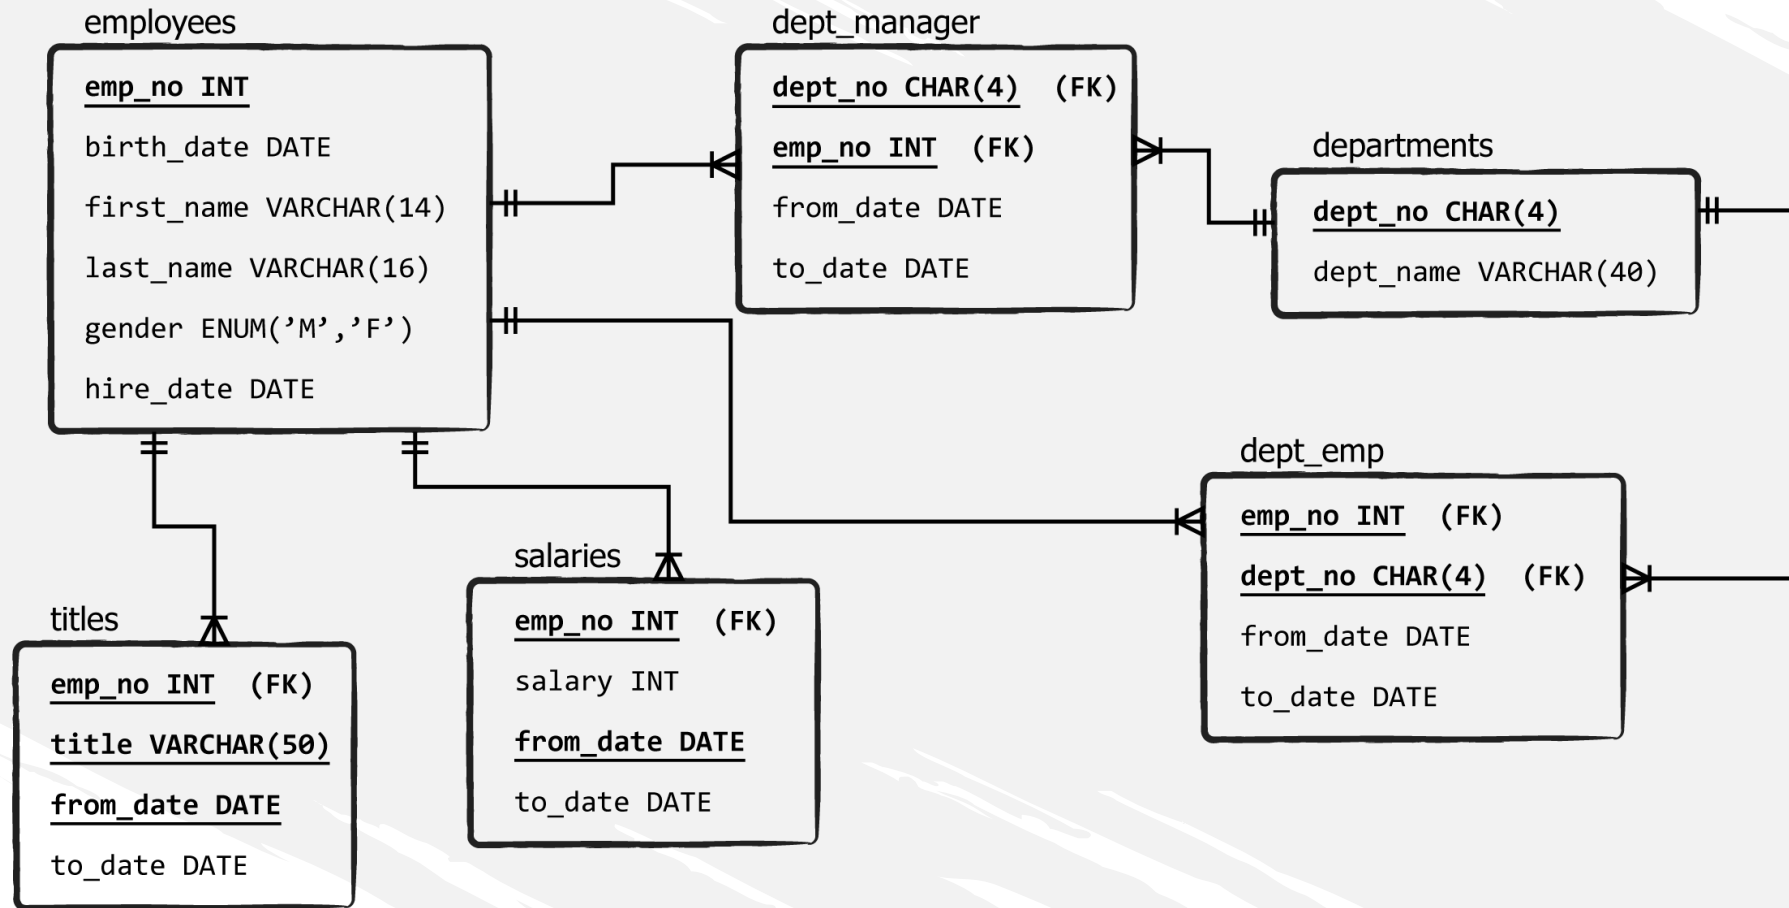
Introduction to Joins

Introduction to Joins


- joins

the SQL tool that allow us to construct a relationship between objects

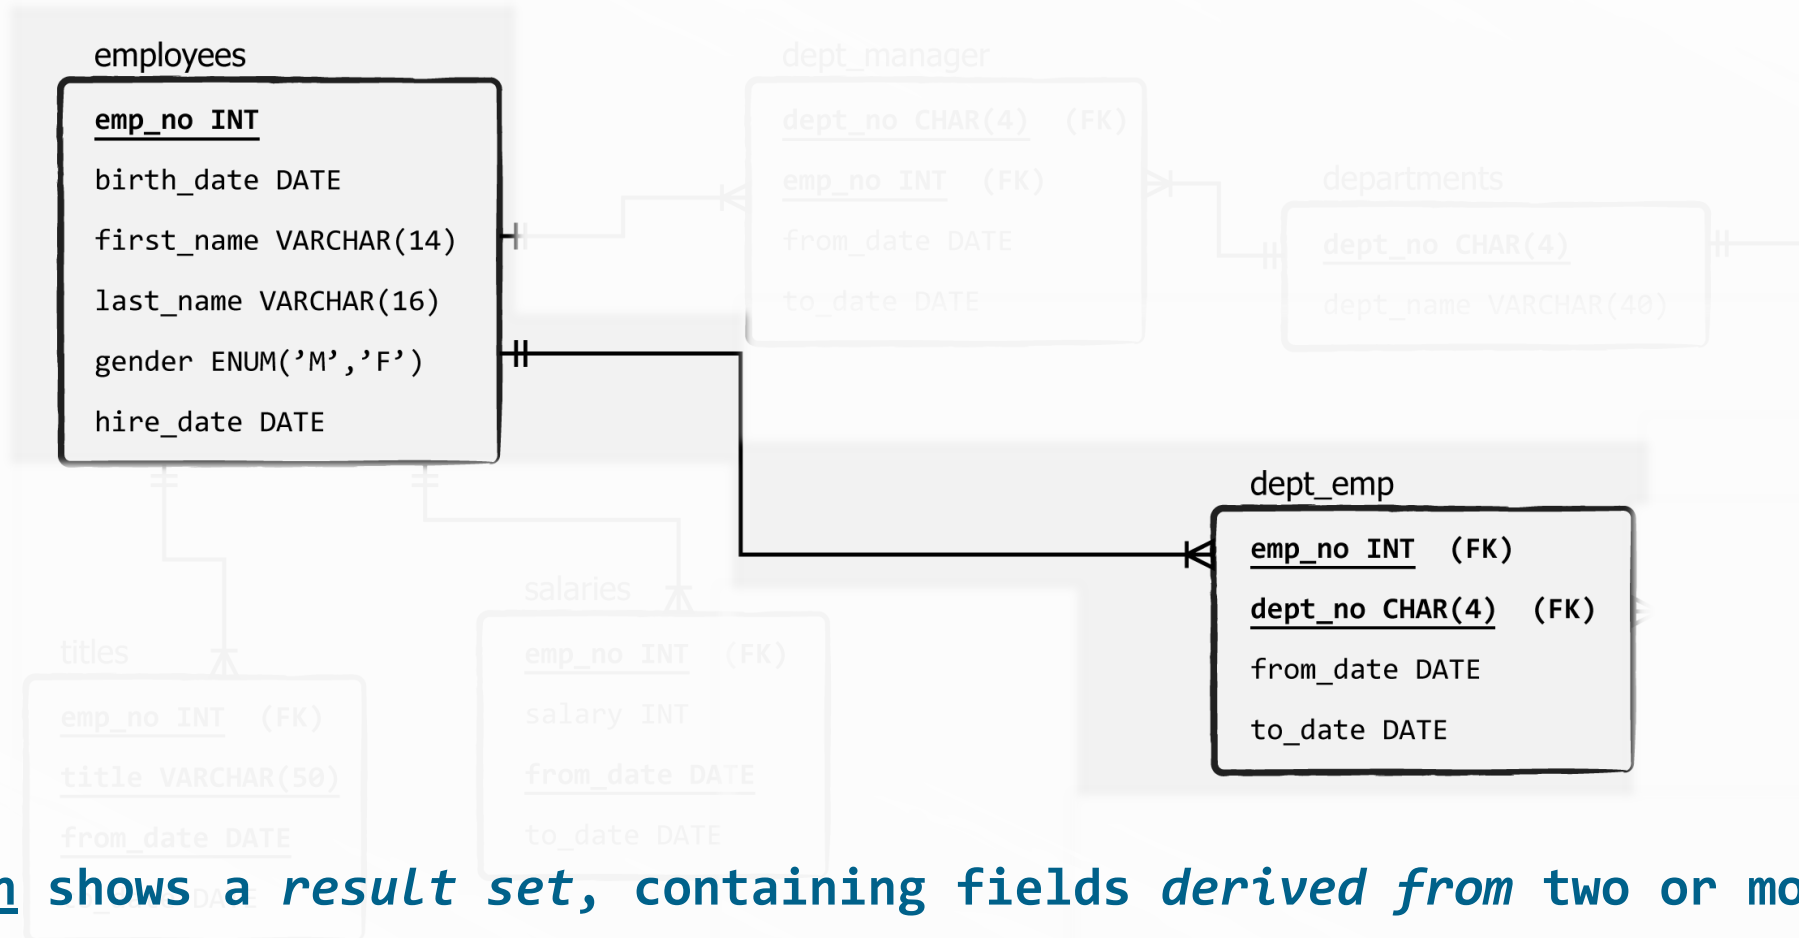
Introduction to Joins



Introduction to Joins



Introduction to Joins



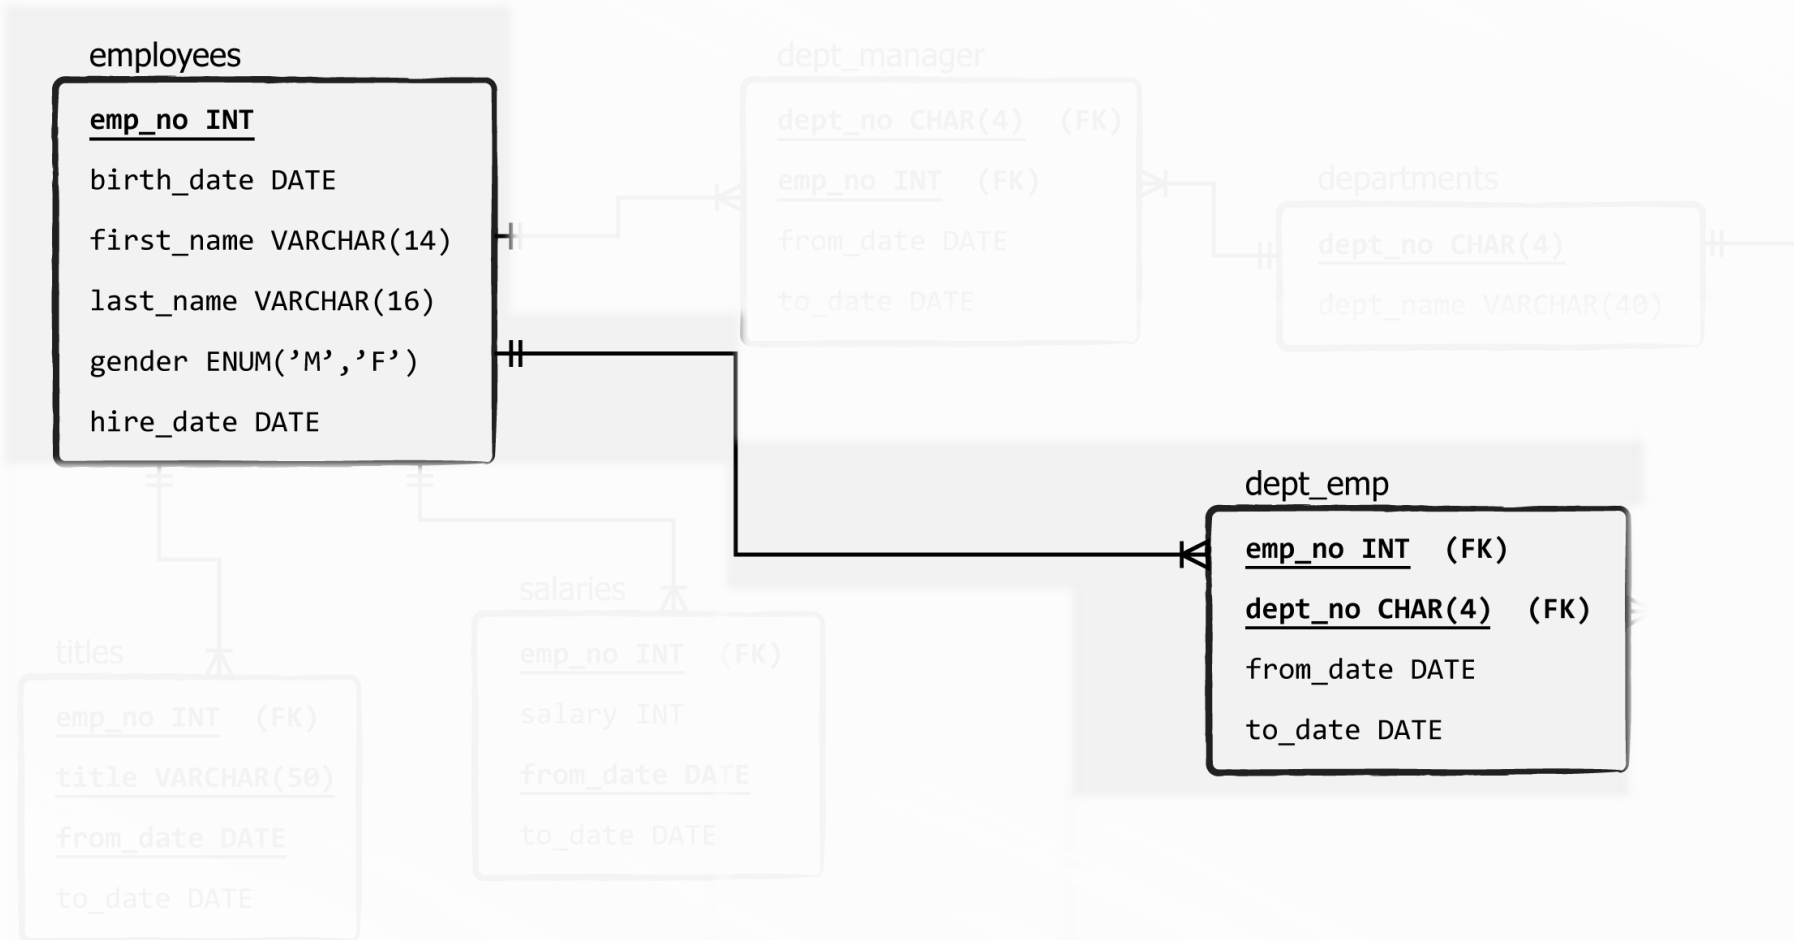
a join shows a *result set*, containing fields *derived from* two or more tables

Introduction to Joins

● joins

- we must find a *related column* from the two tables that contains the same *type* of data

Introduction to Joins

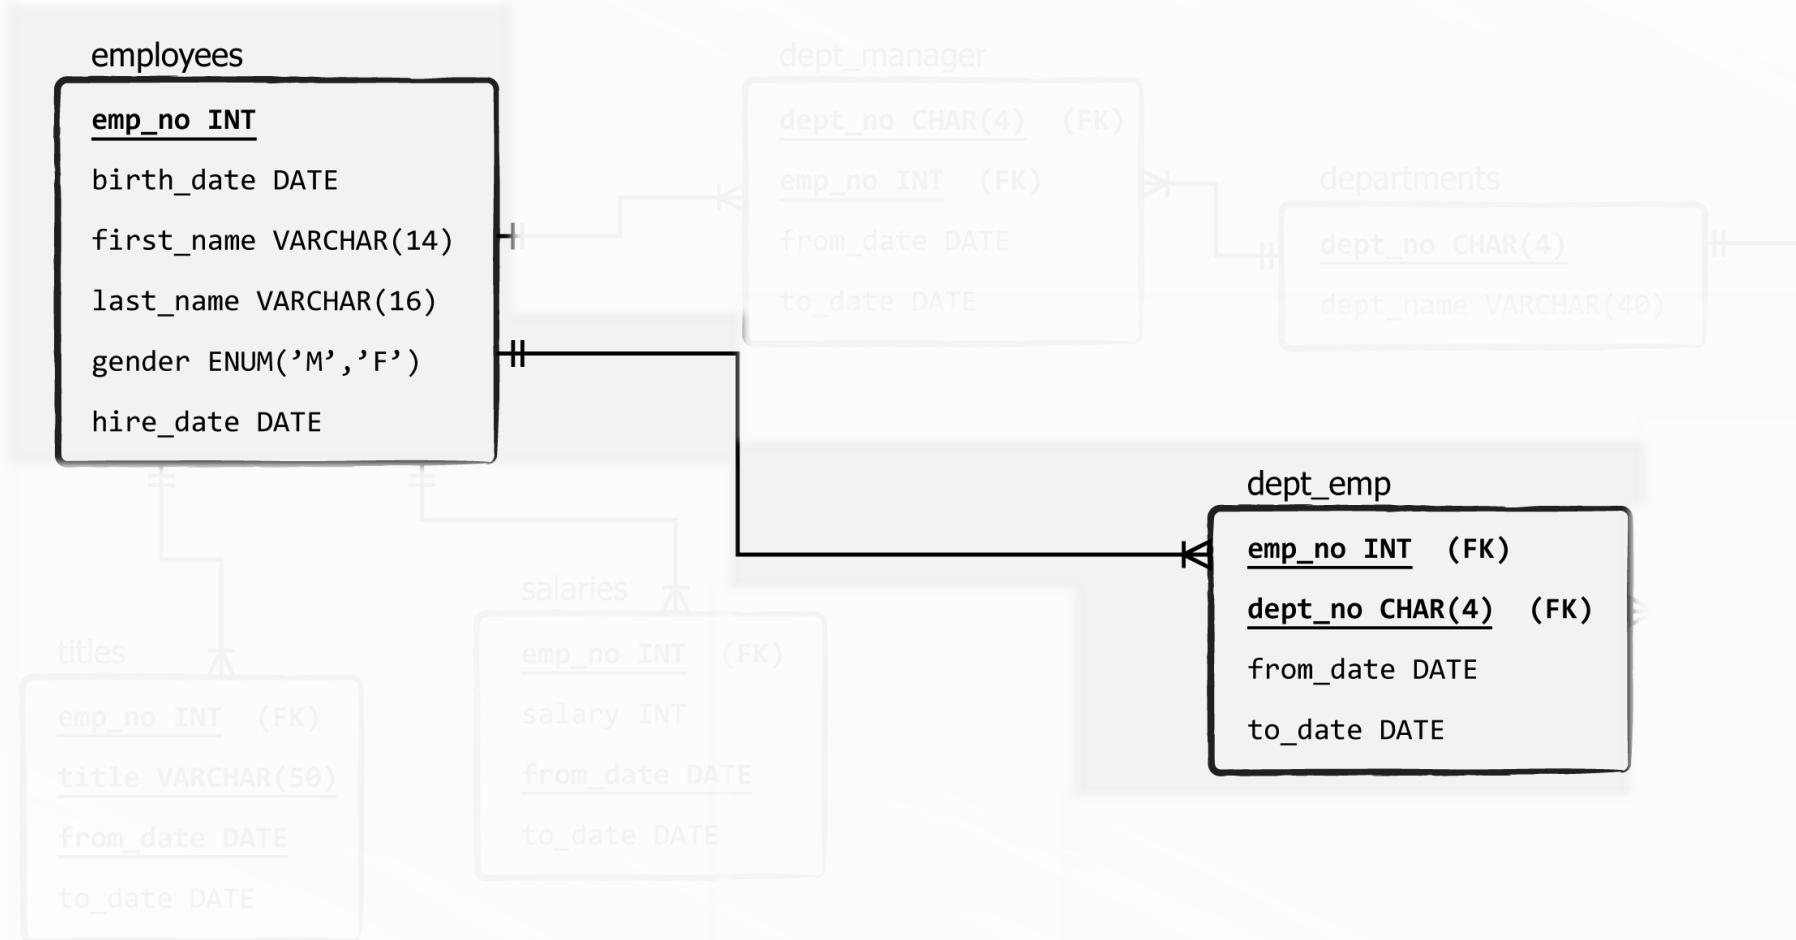


Introduction to Joins

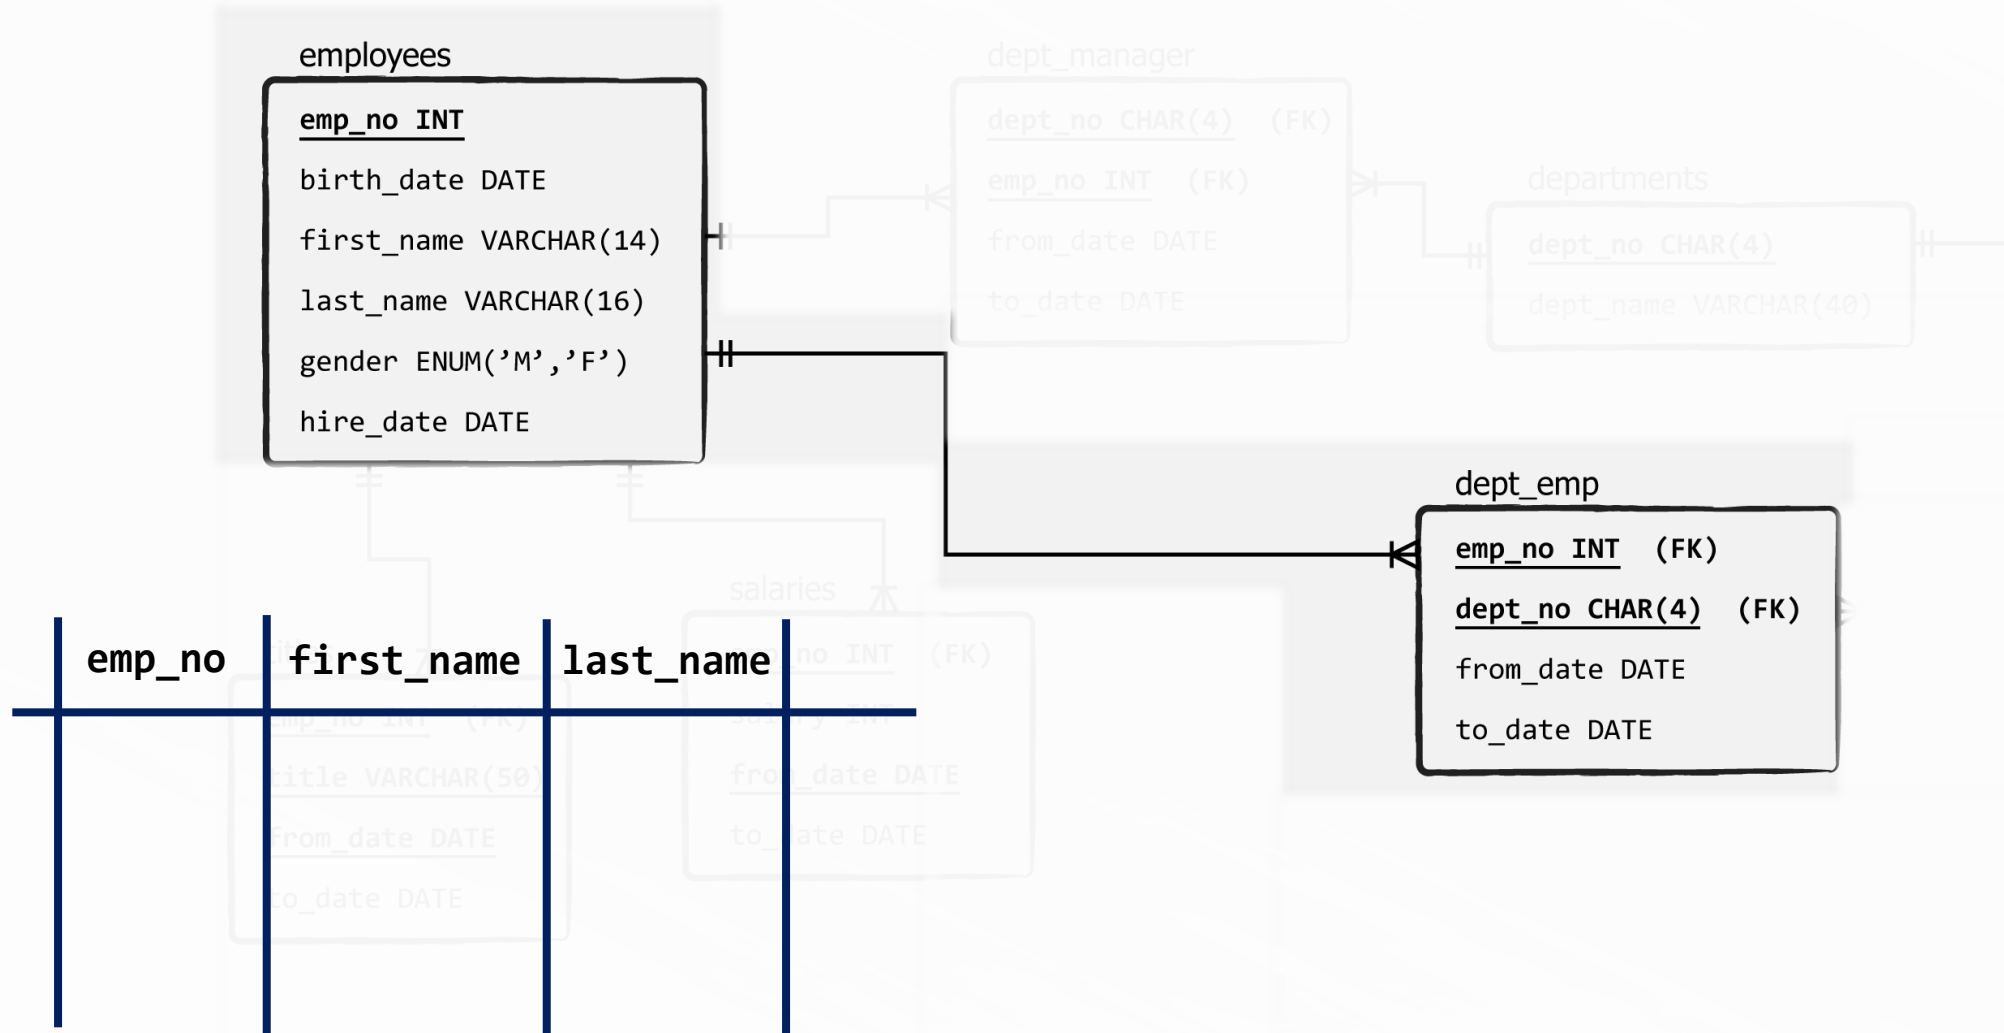
● joins

- we must find a *related column* from the two tables that contains the same *type* of data
- we will be free to add columns from these two tables to our output

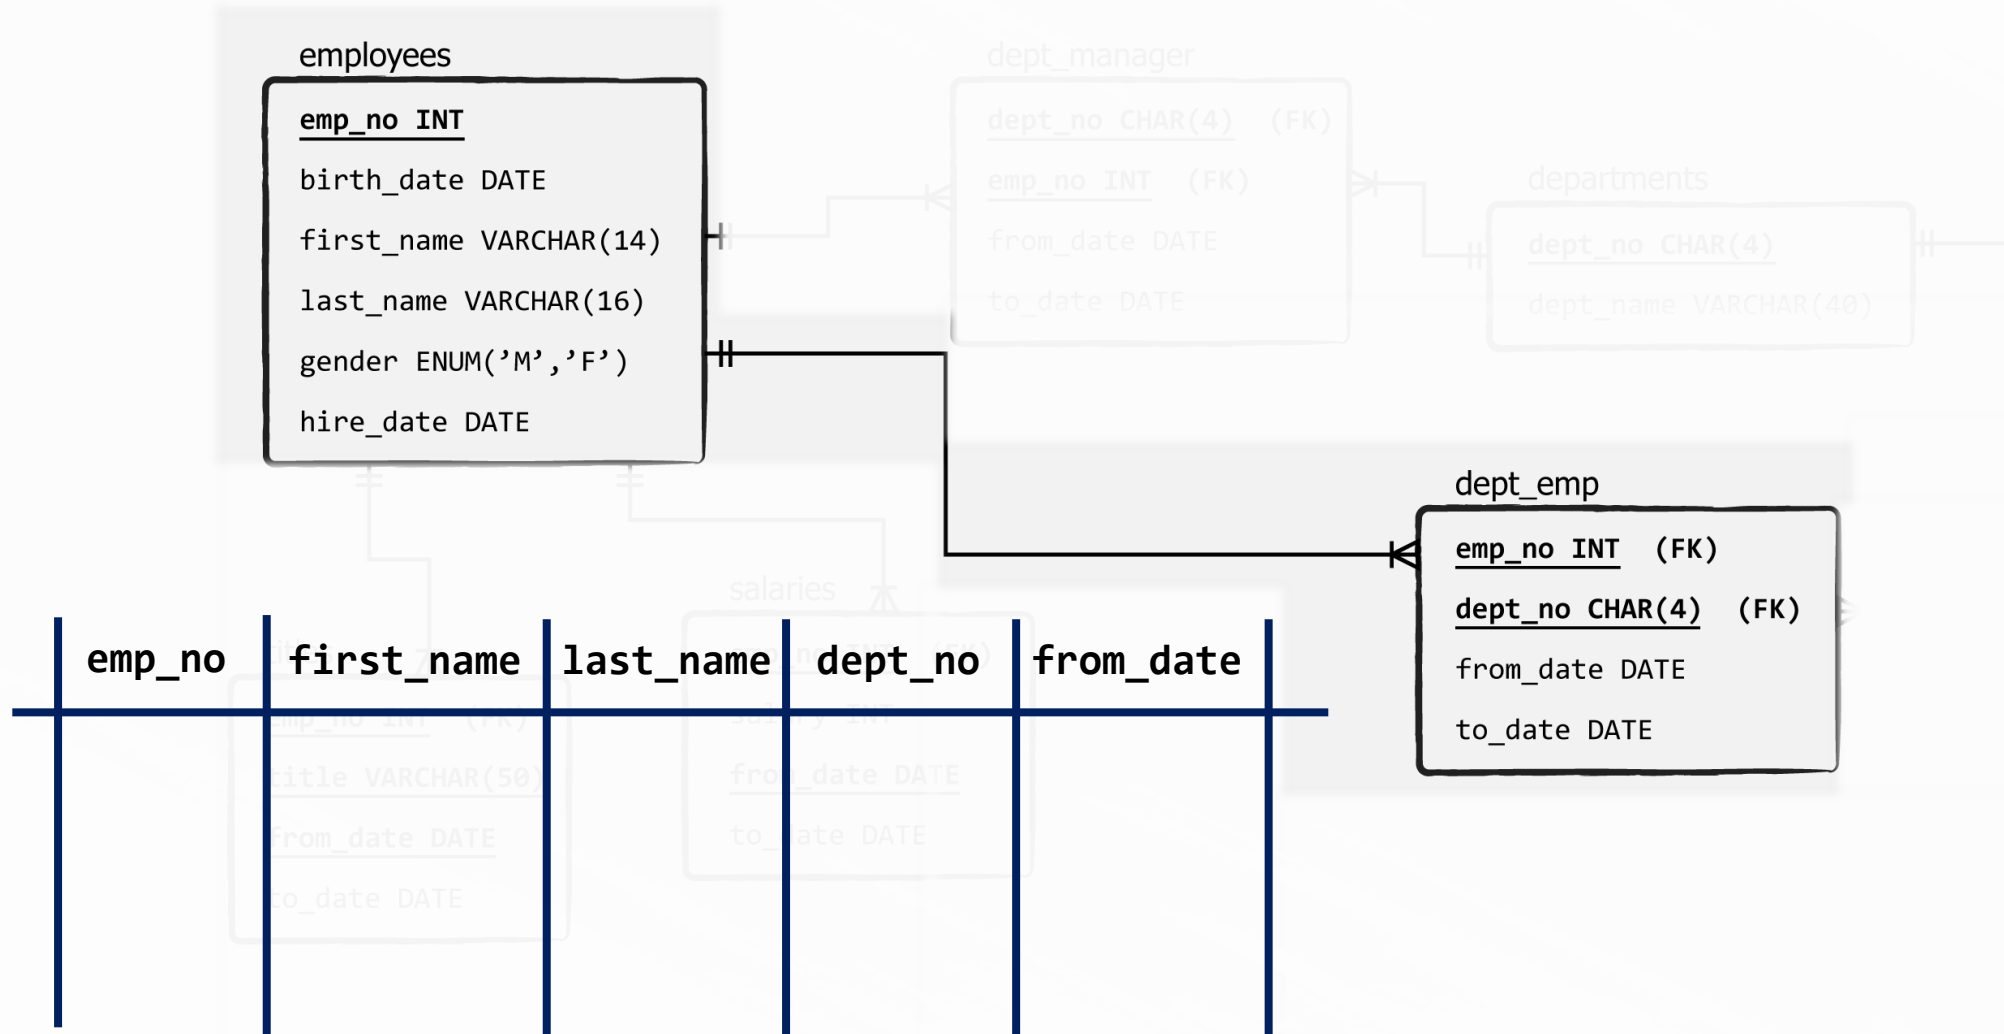
Introduction to Joins




Introduction to Joins



Introduction to Joins



Introduction to Joins

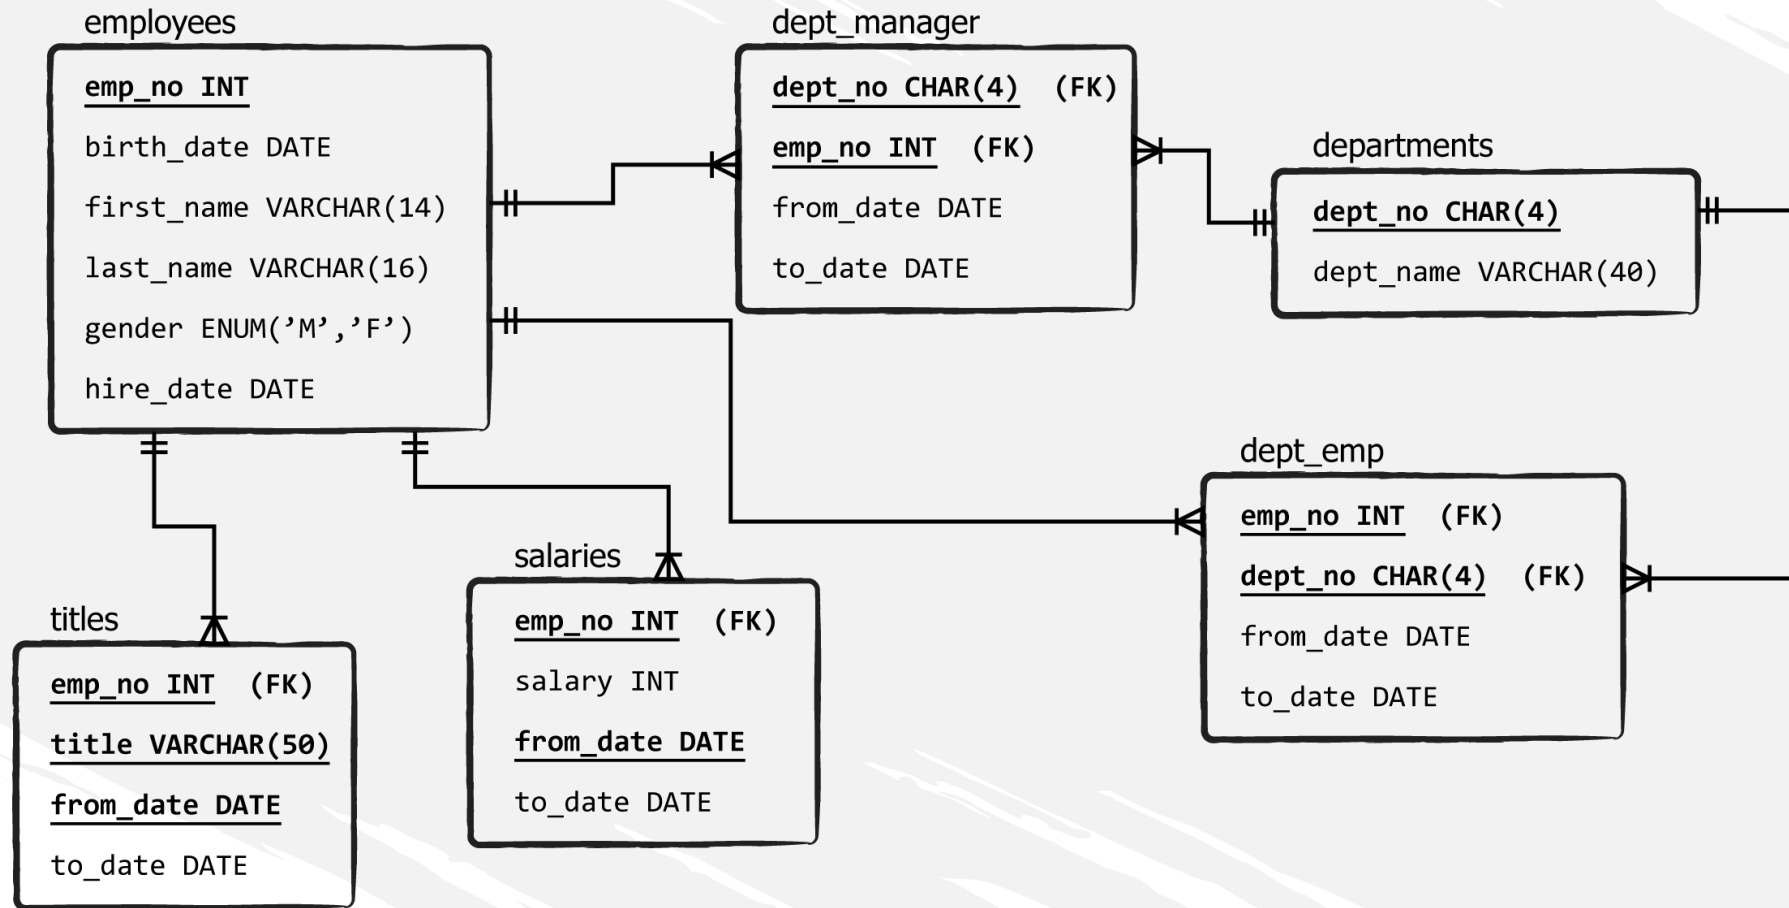


Introduction to Joins

● joins

- the columns you use to relate tables must represent the same object, *such as id*
- the tables you are considering need not be logically adjacent

Introduction to Joins



Introduction to Joins

- We will use two duplicate tables:

- 'departments_dup'
- 'dept_manager_dup'

INNER JOIN


INNER JOIN



INNER JOIN


INNER JOIN

INNER JOIN



INNER JOIN


INNER JOIN



Venn diagram

INNER JOIN

INNER JOIN



Venn diagram

a mathematical tool representing all possible logical relations between a finite collection of sets

INNER JOIN


dept_manager_dup

dept_no CHAR(4)

emp_no INT

from_date DATE

to_date DATE



Venn diagram

a mathematical tool representing all possible logical relations between a finite collection of sets

INNER JOIN


dept_manager_dup

dept_no CHAR(4)

emp_no INT

from_date DATE

to_date DATE



departments_dup

dept_no CHAR(4)

dept_name VARCHAR(40)

Venn diagram

a mathematical tool representing all possible logical relations between a finite collection of sets

INNER JOIN


dept_manager_dup

dept_no CHAR(4)

emp_no INT

from_date DATE

to_date DATE



departments_dup

dept_no CHAR(4)

dept_name VARCHAR(40)

Which will be the related column here?

INNER JOIN


dept_manager_dup

dept_no CHAR(4)

emp_no INT

from_date DATE

to_date DATE



departments_dup

dept_no CHAR(4)

dept_name VARCHAR(40)

Related column: dept_no

INNER JOIN


dept_manager_dup

dept_no CHAR(4)

emp_no INT

from_date DATE

to_date DATE



departments_dup

dept_no CHAR(4)

dept_name VARCHAR(40)

the area that belongs to both circles, which is filled with red, represents all records belonging to *both* the “Department Manager Duplicate” and the “Departments Duplicate” tables

INNER JOIN


dept_manager_dup

dept_no CHAR(4)

emp_no INT

from_date DATE

to_date DATE



departments_dup

dept_no CHAR(4)

dept_name VARCHAR(40)


result set

the area that belongs to both circles, which is filled with red, represents all records belonging to *both* the “Department Manager Duplicate” and the “Departments Duplicate” tables

INNER JOIN

- INNER JOIN

can help us extract this result set




result set

INNER JOIN

- INNER JOIN

can help us extract this result set



result set

INNER JOIN

dept_manager_dup

dept_no CHAR(4)

emp_no INT

from_date DATE

to_date DATE

dept_no CHAR(4)




departments_dup

dept_no CHAR(4)

dept_name VARCHAR(40)

INNER JOIN



matching values = matching records

INNER JOIN

dept_manager_dup


dept_no CHAR(4)

emp_no INT

from_date DATE

to_date DATE

dept_no CHAR(4)



departments_dup

dept_no CHAR(4)

dept_name VARCHAR(40)

matching values = matching records

INNER JOIN


dept_manager_dup

dept_no CHAR(4)

emp_no INT

from_date DATE

to_date DATE




departments_dup

dept_no CHAR(4)

dept_name VARCHAR(40)

INNER JOIN



non-matching values = non-matching records

INNER JOIN


dept_manager_dup

dept_no CHAR(4)

emp_no INT

from_date DATE

to_date DATE



departments_dup

dept_no CHAR(4)

dept_name VARCHAR(40)

non-matching values = non-matching records

INNER JOIN

INNER JOIN



SQL

```
SELECT
    table_1.column_name(s), table_2.column_name(s)
FROM
    table_1
JOIN
    table_2 ON table_1.column_name = table_2.column_name;
```


INNER JOIN

INNER JOIN



SQL

```
SELECT
    t1.column_name, t1.column_name, ..., t2.column_name, ...
FROM
    table_1 t1
JOIN
    table_2 t2 ON t1.column_name = t2.column_name;
```

INNER JOIN


dept_manager_dup

dept_no CHAR(4)

emp_no INT

from_date DATE

to_date DATE



departments_dup

dept_no CHAR(4)

dept_name VARCHAR(40)

INNER JOIN

M


dept_manager_dup

dept_no CHAR(4)

emp_no INT

from_date DATE

to_date DATE



D

departments_dup

dept_no CHAR(4)

dept_name VARCHAR(40)

INNER JOIN


- inner joins extract only records in which the values in the related columns match. Null values, or values appearing in just one of the two tables and not appearing in the other, are not displayed
 - only non-null matching values are in play

INNER JOIN

- And what if such *matching values* did not exist?


INNER JOIN

- And what if such *matching values* did not exist?



INNER JOIN

- And what if such *matching values* did not exist?



Simply, the result set will be empty. There will be no link between the two tables.

Duplicate Records

Duplicate Records

- duplicate records, also known as duplicate rows, are identical rows in an SQL table

Duplicate Records

- duplicate records, also known as duplicate rows, are identical rows in an SQL table
 - for a pair of duplicate records, the values in each column coincide

Duplicate Records

- duplicate rows are not always allowed in a database or a data table

Duplicate Records

- duplicate rows are not always allowed in a database or a data table
 - they are sometimes encountered, especially in new, raw, or uncontrolled data

Duplicate Records

- duplicate rows are not always allowed in a database or a data table
 - they are sometimes encountered, especially in new, raw, or uncontrolled data

here's how to handle duplicates:

GROUP BY the field that differs most among records!

A modern office interior with large windows and a long conference table. The room is brightly lit, and the windows offer a view of a city skyline. The conference table is long and dark, with several chairs arranged around it. The text "LEFT JOIN" is overlaid in the center of the image.

LEFT JOIN


LEFT JOIN



LEFT JOIN

LEFT JOIN

LEFT JOIN



LEFT JOIN


dept_manager_dup

dept_no CHAR(4)

emp_no INT

from_date DATE

to_date DATE



departments_dup

dept_no CHAR(4)

dept_name VARCHAR(40)

LEFT JOIN


dept_manager_dup

dept_no CHAR(4)

emp_no INT

from_date DATE

to_date DATE



departments_dup

dept_no CHAR(4)

dept_name VARCHAR(40)

LEFT JOIN


dept_manager_dup

dept_no CHAR(4)

emp_no INT

from_date DATE

to_date DATE



departments_dup

dept_no CHAR(4)

dept_name VARCHAR(40)

- 1) all matching values of the two tables +
- 2) all values from the left table that match no values from the right table

LEFT JOIN


dept_manager_dup

dept_no CHAR(4)

emp_no INT

from_date DATE

to_date DATE



departments_dup

dept_no CHAR(4)

dept_name VARCHAR(40)

all matching values of the two tables +
all values from the left table that match no values from the right table

LEFT JOIN


dept_manager_dup

dept_no CHAR(4)

emp_no INT

from_date DATE

to_date DATE



departments_dup

dept_no CHAR(4)

dept_name VARCHAR(40)

~~all matching values of the two tables~~ +
all values from the left table that match no values from the right table

LEFT JOIN

dept_manager_dup

dept_no CHAR(4)

emp_no INT

from_date DATE

to_date DATE



departments_dup

dept_no CHAR(4)

dept_name VARCHAR(40)

~~all matching values of the two tables~~ +
all values from the left table that match no values from the right table

LEFT JOIN

● LEFT JOIN

when working with left joins, the order in which you join tables *matters*

LEFT JOIN


dept_manager_dup

dept_no CHAR(4)

emp_no INT

from_date DATE

to_date DATE



departments_dup

dept_no CHAR(4)


dept_name VARCHAR(40)

LEFT JOIN

departments_dup

dept_no CHAR(4)

dept_name VARCHAR(40)



dept_manager_dup

dept_no CHAR(4)

emp_no INT

from_date DATE

to_date DATE


departments_dup

dept_no CHAR(4)

dept_name VARCHAR(40)

dept_no	dept_name
NULL	Public Relations
d001	Marketing
d003	Human Resources
d004	Production
d005	Development
d006	Quality Management
d007	Sales
d008	Research
d009	Customer Service
d010	NULL
d011	NULL

LEFT JOIN



dept_manager_dup

dept_no CHAR(4)


emp_no INT

from_date DATE

to_date DATE


emp_no	dept_no	from_date	to_date
999904	NULL	2017-01-01	NULL
999905	NULL	2017-01-01	NULL
999906	NULL	2017-01-01	NULL
999907	NULL	2017-01-01	NULL
110085	d002	1985-01-01	1989-12-17
110114	d002	1989-12-17	9999-01-01
110183	d003	1985-01-01	1992-03-21
110228	d003	1992-03-21	9999-01-01
110303	d004	1985-01-01	1988-09-09
110344	d004	1988-09-09	1992-08-02
110386	d004	1992-08-02	1996-08-30
110420	d004	1996-08-30	9999-01-01
110511	d005	1985-01-01	1992-04-25
110567	d005	1992-04-25	9999-01-01
110725	d006	1985-01-01	1989-05-06
110765	d006	1989-05-06	1991-09-12
110800	d006	1991-09-12	1994-06-28
110854	d006	1994-06-28	9999-01-01
111035	d007	1985-01-01	1991-03-07
111133	d007	1991-03-07	9999-01-01
111400	d008	1985-01-01	1991-04-08
111534	d008	1991-04-08	9999-01-01
111692	d009	1985-01-01	1988-10-17
111784	d009	1988-10-17	1992-09-08
111877	d009	1992-09-08	1996-01-03
111939	d009	1996-01-03	9999-01-01

LEFT JOIN



INNER join


LEFT JOIN



INNER join


the result set is in the *inner* part of the Venn diagram

LEFT JOIN



LEFT join


LEFT JOIN



LEFT join

in the output obtained you have data from the *outer* part of the Venn diagram too


LEFT JOIN



LEFT join = LEFT OUTER join

in the output obtained you have data from the
outer part of the Venn diagram too

LEFT JOIN



LEFT join = LEFT OUTER join


if you are using a left join, it will always be an
OUTER type of join

LEFT JOIN

- left joins can deliver a list with all records from the left table that do not match any rows from the right table

LEFT JOIN

- left joins can deliver a list with all records from the left table that do not match any rows from the right table



LEFT JOIN

left joins can deliver a list with all records from the left table that do not match any rows from the right table


dept_manager_dup

dept_no CHAR(4)

emp_no INT

from_date DATE

to_date DATE



departments_dup


dept_no CHAR(4)

dept_name VARCHAR(40)

LEFT JOIN

dept_manager_dup

```
dept_no CHAR(4)
emp_no INT
from_date DATE
to_date DATE
```



departments_dup

```
dept_no CHAR(4)
dept_name VARCHAR(40)
```



SQL

```
SELECT
    t1.column_name, t1.column_name, ..., t2.column_name, ...
FROM
    table_1 t1
JOIN
    table_2 t2 ON t1.column_name = t2.column_name
WHERE
    column_name ... IS NULL;
```

A modern conference room with large windows and a long table. The room is empty, with several chairs arranged around the table. The windows offer a view of a city skyline. The text "RIGHT JOIN" is overlaid in the center.

RIGHT JOIN

RIGHT JOIN



RIGHT JOIN

RIGHT JOIN

● RIGHT JOIN

their functionality is identical to LEFT JOINs, with the only difference being that the direction of the operation is inverted

RIGHT JOIN

- RIGHT JOIN = RIGHT OUTER JOIN

their functionality is identical to LEFT JOINs, with the only difference being that the direction of the operation is inverted

RIGHT JOIN


dept_manager_dup

dept_no CHAR(4)

emp_no INT

from_date DATE

to_date DATE



departments_dup

dept_no CHAR(4)

dept_name VARCHAR(40)

RIGHT JOIN


dept_manager_dup

dept_no CHAR(4)

emp_no INT

from_date DATE

to_date DATE



departments_dup

dept_no CHAR(4)

dept_name VARCHAR(40)

RIGHT JOIN


dept_manager_dup

dept_no CHAR(4)

emp_no INT

from_date DATE

to_date DATE



departments_dup

dept_no CHAR(4)

dept_name VARCHAR(40)

RIGHT JOIN


dept_manager_dup

dept_no CHAR(4)

emp_no INT

from_date DATE

to_date DATE



departments_dup

dept_no CHAR(4)

dept_name VARCHAR(40)

dept_manager_dup

dept_no CHAR(4)


emp_no INT

from_date DATE

to_date DATE

emp_no	dept_no	from_date	to_date
999904	NULL	2017-01-01	NULL
999905	NULL	2017-01-01	NULL
999906	NULL	2017-01-01	NULL
999907	NULL	2017-01-01	NULL
110085	d002	1985-01-01	1989-12-17
110114	d002	1989-12-17	9999-01-01
110183	d003	1985-01-01	1992-03-21
110228	d003	1992-03-21	9999-01-01
110303	d004	1985-01-01	1988-09-09
110344	d004	1988-09-09	1992-08-02
110386	d004	1992-08-02	1996-08-30
110420	d004	1996-08-30	9999-01-01
110511	d005	1985-01-01	1992-04-25
110567	d005	1992-04-25	9999-01-01
110725	d006	1985-01-01	1989-05-06
110765	d006	1989-05-06	1991-09-12
110800	d006	1991-09-12	1994-06-28
110854	d006	1994-06-28	9999-01-01
111035	d007	1985-01-01	1991-03-07
111133	d007	1991-03-07	9999-01-01
111400	d008	1985-01-01	1991-04-08
111534	d008	1991-04-08	9999-01-01
111692	d009	1985-01-01	1988-10-17
111784	d009	1988-10-17	1992-09-08
111877	d009	1992-09-08	1996-01-03
111939	d009	1996-01-03	9999-01-01

RIGHT JOIN



departments_dup

dept_no CHAR(4)

dept_name VARCHAR(40)

dept_no	dept_name
NULL	Public Relations
d001	Marketing
d003	Human Resources
d004	Production
d005	Development
d006	Quality Management
d007	Sales
d008	Research
d009	Customer Service
d010	NULL
d011	NULL

RIGHT JOIN

- Whether we run a RIGHT JOIN or a LEFT JOIN with an *inverted* tables order, we will obtain the same output, right?

RIGHT JOIN

- Whether we run a RIGHT JOIN or a LEFT JOIN with an *inverted* tables order, we will obtain the same output, right?

Yes, we will!

dept_manager_dup

dept_no CHAR(4)


emp_no INT

from_date DATE

to_date DATE

emp_no	dept_no	from_date	to_date
999904	NULL	2017-01-01	NULL
999905	NULL	2017-01-01	NULL
999906	NULL	2017-01-01	NULL
999907	NULL	2017-01-01	NULL
110085	d002	1985-01-01	1989-12-17
110114	d002	1989-12-17	9999-01-01
110183	d003	1985-01-01	1992-03-21
110228	d003	1992-03-21	9999-01-01
110303	d004	1985-01-01	1988-09-09
110344	d004	1988-09-09	1992-08-02
110386	d004	1992-08-02	1996-08-30
110420	d004	1996-08-30	9999-01-01
110511	d005	1985-01-01	1992-04-25
110567	d005	1992-04-25	9999-01-01
110725	d006	1985-01-01	1989-05-06
110765	d006	1989-05-06	1991-09-12
110800	d006	1991-09-12	1994-06-28
110854	d006	1994-06-28	9999-01-01
111035	d007	1985-01-01	1991-03-07
111133	d007	1991-03-07	9999-01-01
111400	d008	1985-01-01	1991-04-08
111534	d008	1991-04-08	9999-01-01
111692	d009	1985-01-01	1988-10-17
111784	d009	1988-10-17	1992-09-08
111877	d009	1992-09-08	1996-01-03
111939	d009	1996-01-03	9999-01-01

RIGHT JOIN



departments_dup

dept_no CHAR(4)

dept_name VARCHAR(40)

dept_no	dept_name
NULL	Public Relations
d001	Marketing
d003	Human Resources
d004	Production
d005	Development
d006	Quality Management
d007	Sales
d008	Research
d009	Customer Service
d010	NULL
d011	NULL


departments_dup

dept_no CHAR(4)

dept_name VARCHAR(40)

dept_no	dept_name
NULL	Public Relations
d001	Marketing
d003	Human Resources
d004	Production
d005	Development
d006	Quality Management
d007	Sales
d008	Research
d009	Customer Service
d010	NULL
d011	NULL

RIGHT JOIN



dept_manager_dup

dept_no CHAR(4)

emp_no INT

from_date DATE

to_date DATE

emp_no	dept_no	from_date	to_date
999904	NULL	2017-01-01	NULL
999905	NULL	2017-01-01	NULL
999906	NULL	2017-01-01	NULL
999907	NULL	2017-01-01	NULL
110085	d002	1985-01-01	1989-12-17
110114	d002	1989-12-17	9999-01-01
110183	d003	1985-01-01	1992-03-21
110228	d003	1992-03-21	9999-01-01
110303	d004	1985-01-01	1988-09-09
110344	d004	1988-09-09	1992-08-02
110386	d004	1992-08-02	1996-08-30
110420	d004	1996-08-30	9999-01-01
110511	d005	1985-01-01	1992-04-25
110567	d005	1992-04-25	9999-01-01
110725	d006	1985-01-01	1989-05-06
110765	d006	1989-05-06	1991-09-12
110800	d006	1991-09-12	1994-06-28
110854	d006	1994-06-28	9999-01-01
111035	d007	1985-01-01	1991-03-07
111133	d007	1991-03-07	9999-01-01
111400	d008	1985-01-01	1991-04-08
111534	d008	1991-04-08	9999-01-01
111692	d009	1985-01-01	1988-10-17
111784	d009	1988-10-17	1992-09-08
111877	d009	1992-09-08	1996-01-03
111939	d009	1996-01-03	

dept_manager_dup

dept_no CHAR(4)


emp_no INT

from_date DATE

to_date DATE

emp_no	dept_no	from_date	to_date
999904	NULL	2017-01-01	NULL
999905	NULL	2017-01-01	NULL
999906	NULL	2017-01-01	NULL
999907	NULL	2017-01-01	NULL
110085	d002	1985-01-01	1989-12-17
110114	d002	1989-12-17	9999-01-01
110183	d003	1985-01-01	1992-03-21
110228	d003	1992-03-21	9999-01-01
110303	d004	1985-01-01	1988-09-09
110344	d004	1988-09-09	1992-08-02
110386	d004	1992-08-02	1996-08-30
110420	d004	1996-08-30	9999-01-01
110511	d005	1985-01-01	1992-04-25
110567	d005	1992-04-25	9999-01-01
110725	d006	1985-01-01	1989-05-06
110765	d006	1989-05-06	1991-09-12
110800	d006	1991-09-12	1994-06-28
110854	d006	1994-06-28	9999-01-01
111035	d007	1985-01-01	1991-03-07
111133	d007	1991-03-07	9999-01-01
111400	d008	1985-01-01	1991-04-08
111534	d008	1991-04-08	9999-01-01
111692	d009	1985-01-01	1988-10-17
111784	d009	1988-10-17	1992-09-08
111877	d009	1992-09-08	1996-01-03
111939	d009	1996-01-03	9999-01-01

RIGHT JOIN



departments_dup

dept_no CHAR(4)

dept_name VARCHAR(40)

dept_no	dept_name
NULL	Public Relations
d001	Marketing
d003	Human Resources
d004	Production
d005	Development
d006	Quality Management
d007	Sales
d008	Research
d009	Customer Service
d010	NULL
d011	NULL

RIGHT JOIN

dept_manager_dup


dept_no CHAR(4)

emp_no INT

from_date DATE

to_date DATE

dept_no CHAR(4)



departments_dup

dept_no CHAR(4)

dept_name VARCHAR(40)

RIGHT JOIN

dept_manager_dup


dept_no CHAR(4)

emp_no INT

from_date DATE

to_date DATE

dept_no CHAR(4)



departments_dup

dept_no CHAR(4)

dept_name VARCHAR(40)

matching column = linking column

RIGHT JOIN

dept_manager_dup

dept_no CHAR(4)

emp_no INT

from_date DATE

to_date DATE

dept_no CHAR(4)



departments_dup

dept_no CHAR(4)

dept_name VARCHAR(40)

matching column = linking column

RIGHT JOIN



RIGHT JOIN

RIGHT JOIN


● RIGHT JOIN

when applying a RIGHT JOIN, all the records from the right table will be included in the result set

RIGHT JOIN

RIGHT JOIN


when applying a RIGHT JOIN, all the records from the right table will be included in the result set



RIGHT JOIN

RIGHT JOIN

when applying a RIGHT JOIN, all the records from the right table will be included in the result set




values from the left table will be included only if their *linking column* contains a value coinciding, or *matching*, with a value from the *linking column* of the right table

RIGHT JOIN

RIGHT JOIN

when applying a RIGHT JOIN, all the records from the right table will be included in the result set



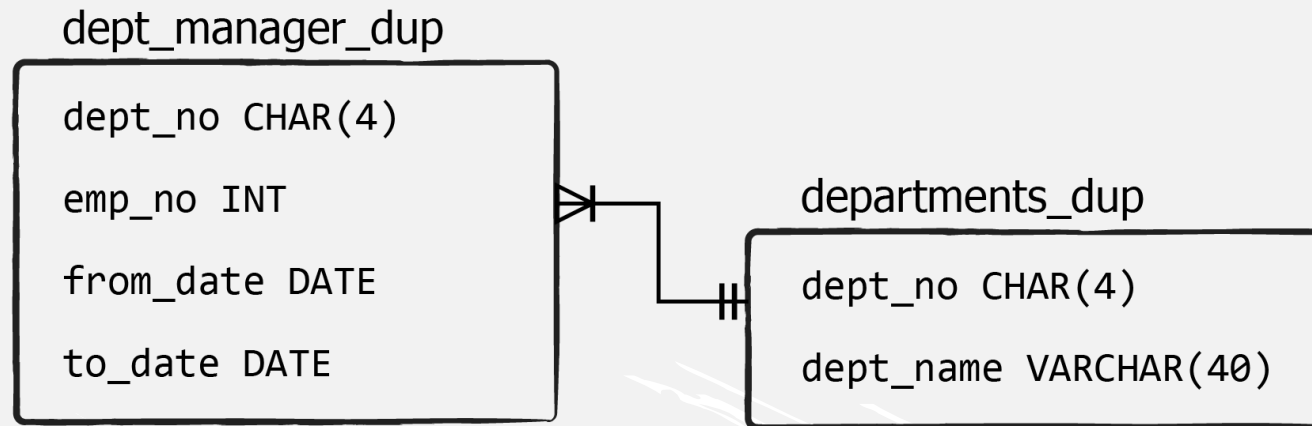
values from the left table will be included only if their *linking column* contains a value coinciding, or *matching*, with a value from the *linking column* of the right table

RIGHT JOIN

- LEFT and RIGHT joins are perfect examples of one-to-many relationships


RIGHT JOIN

- LEFT and RIGHT joins are perfect examples of one-to-many relationships




RIGHT JOIN

- LEFT and RIGHT joins are perfect examples of one-to-many relationships



RIGHT JOIN

- LEFT and RIGHT joins are perfect examples of one-to-many relationships




RIGHT JOIN

- LEFT and RIGHT joins are perfect examples of one-to-many relationships

1 manager

1 department



A modern conference room with large windows and a long table. The room is empty, with several office chairs arranged around the table. The view outside the windows shows a cityscape. The image has a blue tint and a stylized, torn-paper-like border.

The New and the Old Join Syntax

The New and the Old Join Syntax

- Relational Database Essentials
- MySQL Constraints
- SELECT, INSERT, UPDATE, DELETE
- MySQL Aggregate Functions
- INNER JOIN, LEFT JOIN, RIGHT JOIN

The New and the Old Join Syntax

Next:

- Relational Database Essentials
- MySQL Constraints
- SELECT, INSERT, UPDATE, DELETE
- MySQL Aggregate Functions
- INNER JOIN, LEFT JOIN, RIGHT JOIN

The New and the Old Join Syntax

Next:

- Relational Database Essentials
- MySQL Constraints
- SELECT, INSERT, UPDATE, DELETE
- MySQL Aggregate Functions
- INNER JOIN, LEFT JOIN, RIGHT JOIN
- Subqueries, Self-joins
- Stored Routines
- Advanced SQL Tools

The New and the Old Join Syntax




dept_manager_dup

dept_no CHAR(4)

emp_no INT

from_date DATE

to_date DATE




departments_dup

dept_no CHAR(4)

dept_name VARCHAR(40)

The New and the Old Join Syntax




dept_manager_dup

```
dept_no CHAR(4)  
emp_no INT  
from_date DATE  
to_date DATE
```

departments_dup

```
dept_no CHAR(4)  
dept_name VARCHAR(40)
```

The New and the Old Join Syntax



dept_manager_dup

dept_no CHAR(4)

emp_no INT

from_date DATE


to_date DATE

departments_dup

dept_no CHAR(4)

dept_name VARCHAR(40)

The New and the Old Join Syntax



dept_manager_dup

dept_no CHAR(4)

emp_no INT

from_date DATE


to_date DATE

departments_dup

dept_no CHAR(4)

dept_name VARCHAR(40)

The New and the Old Join Syntax



dept_manager_dup

dept_no CHAR(4)

emp_no INT

from_date DATE


to_date DATE

departments_dup

dept_no CHAR(4)

dept_name VARCHAR(40)

The New and the Old Join Syntax



connection points

dept_manager_dup

dept_no CHAR(4)

emp_no INT

from_date DATE

to_date DATE

departments_dup

dept_no CHAR(4)

dept_name VARCHAR(40)

The New and the Old Join Syntax


dept_manager_dup

dept_no CHAR(4)

emp_no INT

from_date DATE

to_date DATE



connection points

departments_dup

dept_no CHAR(4)

dept_name VARCHAR(40)

The New and the Old Join Syntax


emp_no	dept_no	from_date	to_date
999904	NULL	2017-01-01	NULL
999905	NULL	2017-01-01	NULL
999906	NULL	2017-01-01	NULL
999907	NULL	2017-01-01	NULL
110085	d002	1985-01-01	1989-12-17
110114	d002	1989-12-17	9999-01-01
110183	d003	1985-01-01	1992-03-21
110228	d003	1992-03-21	9999-01-01
110303	d004	1985-01-01	1988-09-09
110344	d004	1988-09-09	1992-08-02
110386	d004	1992-08-02	1996-08-30
110420	d004	1996-08-30	9999-01-01
110511	d005	1985-01-01	1992-04-25
110567	d005	1992-04-25	9999-01-01
110725	d006	1985-01-01	1989-05-06
110765	d006	1989-05-06	1991-09-12
110800	d006	1991-09-12	1994-06-28
110854	d006	1994-06-28	9999-01-01
111035	d007	1985-01-01	1991-03-07
111133	d007	1991-03-07	9999-01-01
111400	d008	1985-01-01	1991-04-08
111534	d008	1991-04-08	9999-01-01
111692	d009	1985-01-01	1988-10-17
111784	d009	1988-10-17	1992-09-08
111877	d009	1992-09-08	1996-01-03
111939	d009	1996-01-03	9999-01-01

dept_manager_dup

```
dept_no CHAR(4)
emp_no INT
from_date DATE
to_date DATE
```

dept_no	emp_no	dept_name
d003	110228	Human Resources
d003	110183	Human Resources
d004	110344	Production
d004	110420	Production
d004	110303	Production
d004	110386	Production
d005	110567	Development
d005	110511	Development
d006	110800	Quality Management
d006	110765	Quality Management
d006	110854	Quality Management
d006	110725	Quality Management
d007	111035	Sales
d007	111133	Sales
d008	111400	Research
d008	111534	Research
d009	111784	Customer Service
d009	111939	Customer Service
d009	111692	Customer Service
d009	111877	Customer Service

dept_no	dept_name
NULL	Public Relations
d001	Marketing
d003	Human Resources
d004	Production
d005	Development
d006	Quality Management
d007	Sales
d008	Research
d009	Customer Service
d010	NULL
d011	NULL



connection points

departments_dup

```
dept_no CHAR(4)
dept_name VARCHAR(40)
```

The New and the Old Join Syntax


emp_no	dept_no		from_date	to_date
999904	NULL	✗	2017-01-01	NULL
999905	NULL	✗	2017-01-01	NULL
999906	NULL	✗	2017-01-01	NULL
999907	NULL	✗	2017-01-01	NULL
110085	d002	✗	1985-01-01	1989-12-17
110114	d002	✗	1989-12-17	9999-01-01
110183	d003	✓	1985-01-01	1992-03-21
110228	d003	✓	1992-03-21	9999-01-01
110303	d004	✓	1985-01-01	1988-09-09
110344	d004	✓	1988-09-09	1992-08-02
110386	d004	✓	1992-08-02	1996-08-30
110420	d004	✓	1996-08-30	9999-01-01
110511	d005	✓	1985-01-01	1992-04-25
110567	d005	✓	1992-04-25	9999-01-01
110725	d006	✓	1985-01-01	1989-05-06
110765	d006	✓	1989-05-06	1991-09-12
110800	d006	✓	1991-09-12	1994-06-28
110854	d006	✓	1994-06-28	9999-01-01
111035	d007	✓	1985-01-01	1991-03-07
111133	d007	✓	1991-03-07	9999-01-01
111400	d008	✓	1985-01-01	1991-04-08
111534	d008	✓	1991-04-08	9999-01-01
111692	d009	✓	1985-01-01	1988-10-17
111784	d009	✓	1988-10-17	1992-09-08
111877	d009	✓	1992-09-08	1996-01-03
111939	d009	✓	1996-01-03	9999-01-01

dept_manager_dup

```
dept_no CHAR(4)
emp_no INT
from_date DATE
to_date DATE
```

	dept_no	emp_no	dept_name
	d003	110228	Human Resources
	d003	110183	Human Resources
	d004	110344	Production
	d004	110420	Production
	d004	110303	Production
	d004	110386	Production
	d005	110567	Development
	d005	110511	Development
	d006	110800	Quality Management
	d006	110765	Quality Management
	d006	110854	Quality Management
	d006	110725	Quality Management
	d007	111035	Sales
	d007	111133	Sales
	d008	111400	Research
	d008	111534	Research
	d009	111784	Customer Service
	d009	111939	Customer Service
	d009	111692	Customer Service
	d009	111877	Customer Service

dept_no	dept_name
NULL	Public Relations
d001	Marketing
d003	Human Resources
d004	Production
d005	Development
d006	Quality Management
d007	Sales
d008	Research
d009	Customer Service
d010	NULL
d011	NULL



connection points

departments_dup

```
dept_no CHAR(4)
dept_name VARCHAR(40)
```

The New and the Old Join Syntax


emp_no	dept_no		from_date	to_date
999904	NULL	✗	2017-01-01	NULL
999905	NULL	✗	2017-01-01	NULL
999906	NULL	✗	2017-01-01	NULL
999907	NULL	✗	2017-01-01	NULL
110085	d002	✗	1985-01-01	1989-12-17
110114	d002	✗	1989-12-17	9999-01-01
110183	d003	✓	1985-01-01	1992-03-21
110228	d003	✓	1992-03-21	9999-01-01
110303	d004	✓	1985-01-01	1988-09-09
110344	d004	✓	1988-09-09	1992-08-02
110386	d004	✓	1992-08-02	1996-08-30
110420	d004	✓	1996-08-30	9999-01-01
110511	d005	✓	1985-01-01	1992-04-25
110567	d005	✓	1992-04-25	9999-01-01
110725	d006	✓	1985-01-01	1989-05-06
110765	d006	✓	1989-05-06	1991-09-12
110800	d006	✓	1991-09-12	1994-06-28
110854	d006	✓	1994-06-28	9999-01-01
111035	d007	✓	1985-01-01	1991-03-07
111133	d007	✓	1991-03-07	9999-01-01
111400	d008	✓	1985-01-01	1991-04-08
111534	d008	✓	1991-04-08	9999-01-01
111692	d009	✓	1985-01-01	1988-10-17
111784	d009	✓	1988-10-17	1992-09-08
111877	d009	✓	1992-09-08	1996-01-03
111939	d009	✓	1996-01-03	9999-01-01

dept_manager_dup

```
dept_no CHAR(4)
emp_no INT
from_date DATE
to_date DATE
```

	dept_no	emp_no	dept_name
	d003	110228	Human Resources
	d003	110183	Human Resources
	d004	110344	Production
	d004	110420	Production
	d004	110303	Production
	d004	110386	Production
	d005	110567	Development
	d005	110511	Development
	d006	110800	Quality Management
	d006	110765	Quality Management
	d006	110854	Quality Management
	d006	110725	Quality Management
	d007	111035	Sales
	d007	111133	Sales
	d008	111400	Research
	d008	111534	Research
	d009	111784	Customer Service
	d009	111939	Customer Service
	d009	111692	Customer Service
	d009	111877	Customer Service

dept_no		dept_name
NULL	✗	Public Relations
d001	✗	Marketing
d003	✓	Human Resources
d004	✓	Production
d005	✓	Development
d006	✓	Quality Management
d007	✓	Sales
d008	✓	Research
d009	✓	Customer Service
d010	✗	NULL
d011	✗	NULL



connection points

departments_dup

```
dept_no CHAR(4)
dept_name VARCHAR(40)
```

The New and the Old Join Syntax


emp_no	dept_no	from_date	to_date
999904	NULL	×	2017-01-01
999905	NULL	×	2017-01-01
999906	NULL	×	2017-01-01
999907	NULL	×	2017-01-01
110085	d002	×	1985-01-01
110114	d002	×	1989-12-17
110183	d003	✓	1985-01-01
110228	d003	✓	1992-03-21
110303	d004	✓	1985-01-01
110344	d004	✓	1988-09-09
110386	d004	✓	1992-08-02
110420	d004	✓	1996-08-30
110511	d005	✓	1985-01-01
110567	d005	✓	1992-04-25
110725	d006	✓	1985-01-01
110765	d006	✓	1989-05-06
110800	d006	✓	1991-09-12
110854	d006	✓	1994-06-28
111035	d007	✓	1985-01-01
111133	d007	✓	1991-03-07
111400	d008	✓	1985-01-01
111534	d008	✓	1991-04-08
111692	d009	✓	1985-01-01
111784	d009	✓	1988-10-17
111877	d009	✓	1992-09-08
111939	d009	✓	1996-01-03

dept_no	emp_no	dept_name
d003	110228	Human Resources
d003	110183	Human Resources
d004	110344	Production
d004	110420	Production
d004	110303	Production
d004	110386	Production
d005	110567	Development
d005	110511	Development
d006	110800	Quality Management
d006	110765	Quality Management
d006	110854	Quality Management
d006	110725	Quality Management
d007	111035	Sales
d007	111133	Sales
d008	111400	Research
d008	111534	Research
d009	111784	Customer Service
d009	111939	Customer Service
d009	111692	Customer Service
d009	111877	Customer Service

dept_no		dept_name
NULL	✗	Public Relations
d001	✗	Marketing
d003	✓	Human Resources
d004	✓	Production
d005	✓	Development
d006	✓	Quality Management
d007	✓	Sales
d008	✓	Research
d009	✓	Customer Service
d010	✗	NULL
d011	✗	NULL

dept_manager_dup

```
dept_no CHAR(4)
emp_no INT
from_date DATE
to_date DATE
```



connection points

departments_dup

```
dept_no CHAR(4)
dept_name VARCHAR(40)
```

The New and the Old Join Syntax

- WHERE (*the Old Join Syntax*)



SQL

```
SELECT
    t1.column_name, t1.column_name, ..., t2.column_name, ...
FROM
    table_1 t1,
    table_2 t2
WHERE
    t1.column_name = t2.column_name;
```


The New and the Old Join Syntax

- JOIN or WHERE?

The New and the Old Join Syntax

- JOIN or WHERE?

- the retrieved output is *identical*

The New and the Old Join Syntax

● JOIN or WHERE?

- the retrieved output is *identical*
- using WHERE is more *time-consuming*

The New and the Old Join Syntax

● JOIN or WHERE?

- the retrieved output is *identical*
- using WHERE is more *time-consuming*
- the WHERE syntax is perceived as *morally old* and is rarely employed by professionals

The New and the Old Join Syntax

● JOIN or WHERE?

- the retrieved output is *identical*
- using WHERE is more *time-consuming*
- the WHERE syntax is perceived as *morally old* and is rarely employed by professionals
- the JOIN syntax allows you to modify the connection between tables easily

The New and the Old Join Syntax

JOIN (*the New Join Syntax*) vs WHERE (*the Old Join Syntax*)

A modern office interior with large windows and a conference table. The room is brightly lit, and the view outside the windows shows a cityscape. The text "JOIN and WHERE Used Together" is overlaid on the image.

JOIN and WHERE Used Together

JOIN and WHERE Used Together

- JOIN + WHERE

JOIN and WHERE Used Together

● JOIN + WHERE

JOIN:

- used for connecting the “employees” and “salaries” tables

JOIN and WHERE Used Together

● JOIN + WHERE


JOIN:

- used for connecting the “employees” and “salaries” tables

WHERE:

- used to define the condition or conditions that will determine which will be the connecting points between the two tables

JOIN and WHERE Used Together



A modern conference room with large windows and a long table. The room is empty, with several chairs arranged around the table. The windows offer a view of a city skyline. The text "CROSS JOIN" is overlaid in the center.

CROSS JOIN

CROSS JOIN

CROSS JOIN

- a cross join will take the values from a certain table and connect them with all the values from the tables we want to join it with

CROSS JOIN

- a cross join will take the values from a certain table and connect them with all the values from the tables we want to join it with

- INNER JOIN

CROSS JOIN

- a cross join will take the values from a certain table and connect them with all the values from the tables we want to join it with

- INNER JOIN

- typically connects *only the matching values*

CROSS JOIN

- a cross join will take the values from a certain table and connect them with all the values from the tables we want to join it with

- INNER JOIN

- typically connects *only the matching values*

- CROSS JOIN

CROSS JOIN

- a cross join will take the values from a certain table and connect them with all the values from the tables we want to join it with

- INNER JOIN

- typically connects *only the matching values*

- CROSS JOIN

- connects *all the values*, not just those that match

CROSS JOIN

- a cross join will take the values from a certain table and connect them with all the values from the tables we want to join it with

- INNER JOIN

- typically connects *only the matching values*

- CROSS JOIN

- connects *all the values*, not just those that match
- the Cartesian product of the values of two or more sets

CROSS JOIN

- a cross join will take the values from a certain table and connect them with all the values from the tables we want to join it with

- INNER JOIN

- typically connects *only the matching values*

- CROSS JOIN

- connects *all the values*, not just those that match
- the Cartesian product of the values of two or more sets
- particularly useful when the tables in a database are not well connected

FRAGILE

HANDLE WITH CARE

JOIN More Than Two Tables in SQL

JOIN More Than Two Tables in SQL

- when creating a query that joins multiple tables, you must back it with *strong intuition* and a *crystal-clear idea* of how you would like the tables to be connected

JOIN More Than Two Tables in SQL

first_name	last_name	

JOIN More Than Two Tables in SQL

first_name	last_name	hire_date	

JOIN More Than Two Tables in SQL

first_name	last_name	hire_date	from_date	

JOIN More Than Two Tables in SQL

first_name	last_name	hire_date	from_date	dept_name

JOIN More Than Two Tables in SQL

first_name	last_name	hire_date	from_date	dept_name

A modern conference room with large windows and a long table. The room is empty, with several chairs arranged around the table. The view outside the windows shows a cityscape. The image has a blue tint and a stylized, torn-paper-like border.

Tips and Tricks for Joins

Tips and Tricks for Joins

dept_name	
-----------	--


Tips and Tricks for Joins

dept_name	average_salary

Tips and Tricks for Joins

dept_name	average_salary

Tips and Tricks for Joins



Tips and Tricks for Joins



JOINS

Tips and Tricks for Joins

● JOINS

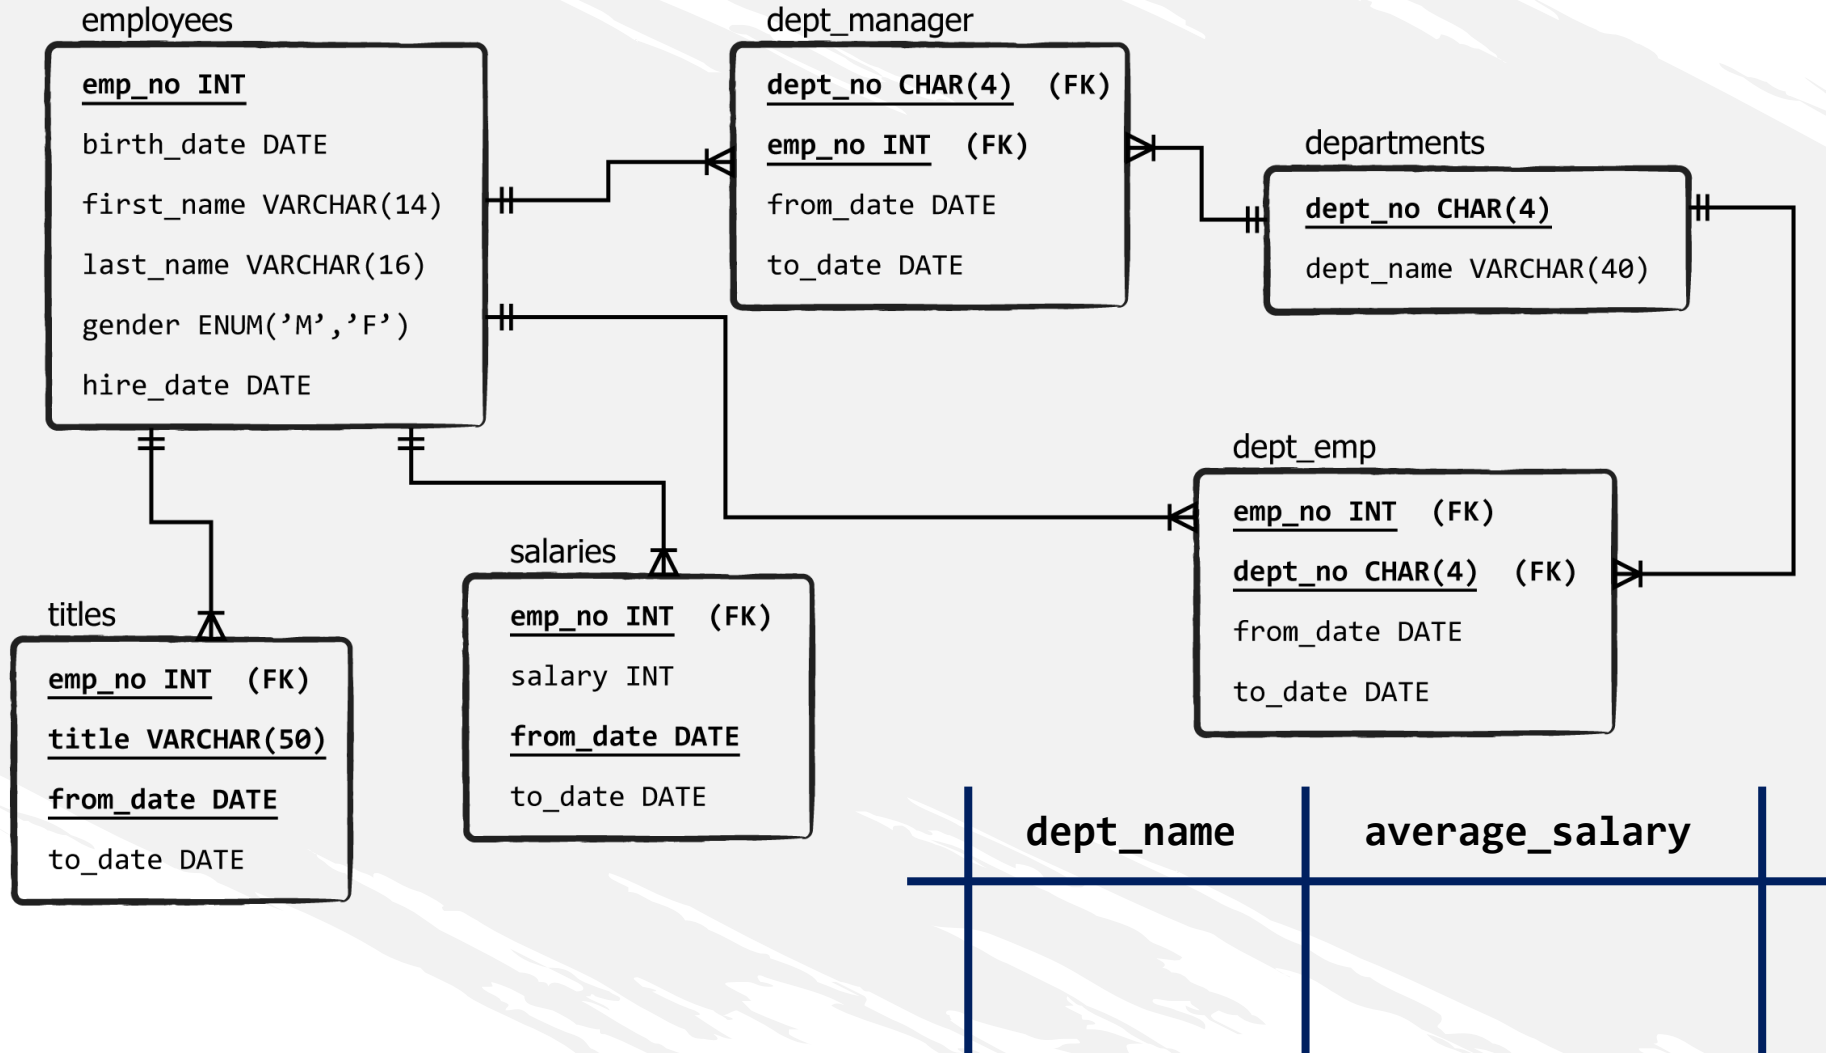
- one should look for key columns, which are common between *the tables involved in the analysis* and are necessary to solve the task at hand

Tips and Tricks for Joins

● JOINS

- one should look for key columns, which are common between *the tables involved in the analysis* and are necessary to solve the task at hand
- these columns do not need to be foreign or private keys

Tips and Tricks for Joins



UNION vs UNION ALL

UNION vs UNION ALL



UNION ALL

UNION vs UNION ALL

- UNION ALL

used to combine a few SELECT statements in a single output

UNION vs UNION ALL

● UNION ALL

used to combine a few SELECT statements in a single output

- you can think of it as a tool that allows you to *unify tables*

UNION vs UNION ALL

- UNION ALL

used to combine a few SELECT statements in a single output



SQL

```
SELECT
    N columns
FROM
    table_1
UNION ALL SELECT
    N columns
FROM
    table_2;
```

UNION vs UNION ALL

- UNION ALL

used to combine a few SELECT statements in a single output



SQL

```
SELECT
    N columns
FROM
    table_1
UNION ALL SELECT
    N columns
FROM
    table_2;
```

We have to select the same number of columns from each table.

UNION vs UNION ALL

● UNION ALL

used to combine a few SELECT statements in a single output



SQL

```
SELECT
    N columns
FROM
    table_1
UNION ALL SELECT
    N columns
FROM
    table_2;
```

We have to select the same number of columns from each table.

These columns should have the same name,

UNION vs UNION ALL

● UNION ALL

used to combine a few SELECT statements in a single output



SQL

```
SELECT
    N columns
FROM
    table_1
UNION ALL SELECT
    N columns
FROM
    table_2;
```

We have to select the same number of columns from each table.

These columns should have the same name, should be in the same order,

UNION vs UNION ALL

● UNION ALL

used to combine a few SELECT statements in a single output




SQL

```
SELECT
    N columns
FROM
    table_1
UNION ALL SELECT
    N columns
FROM
    table_2;
```

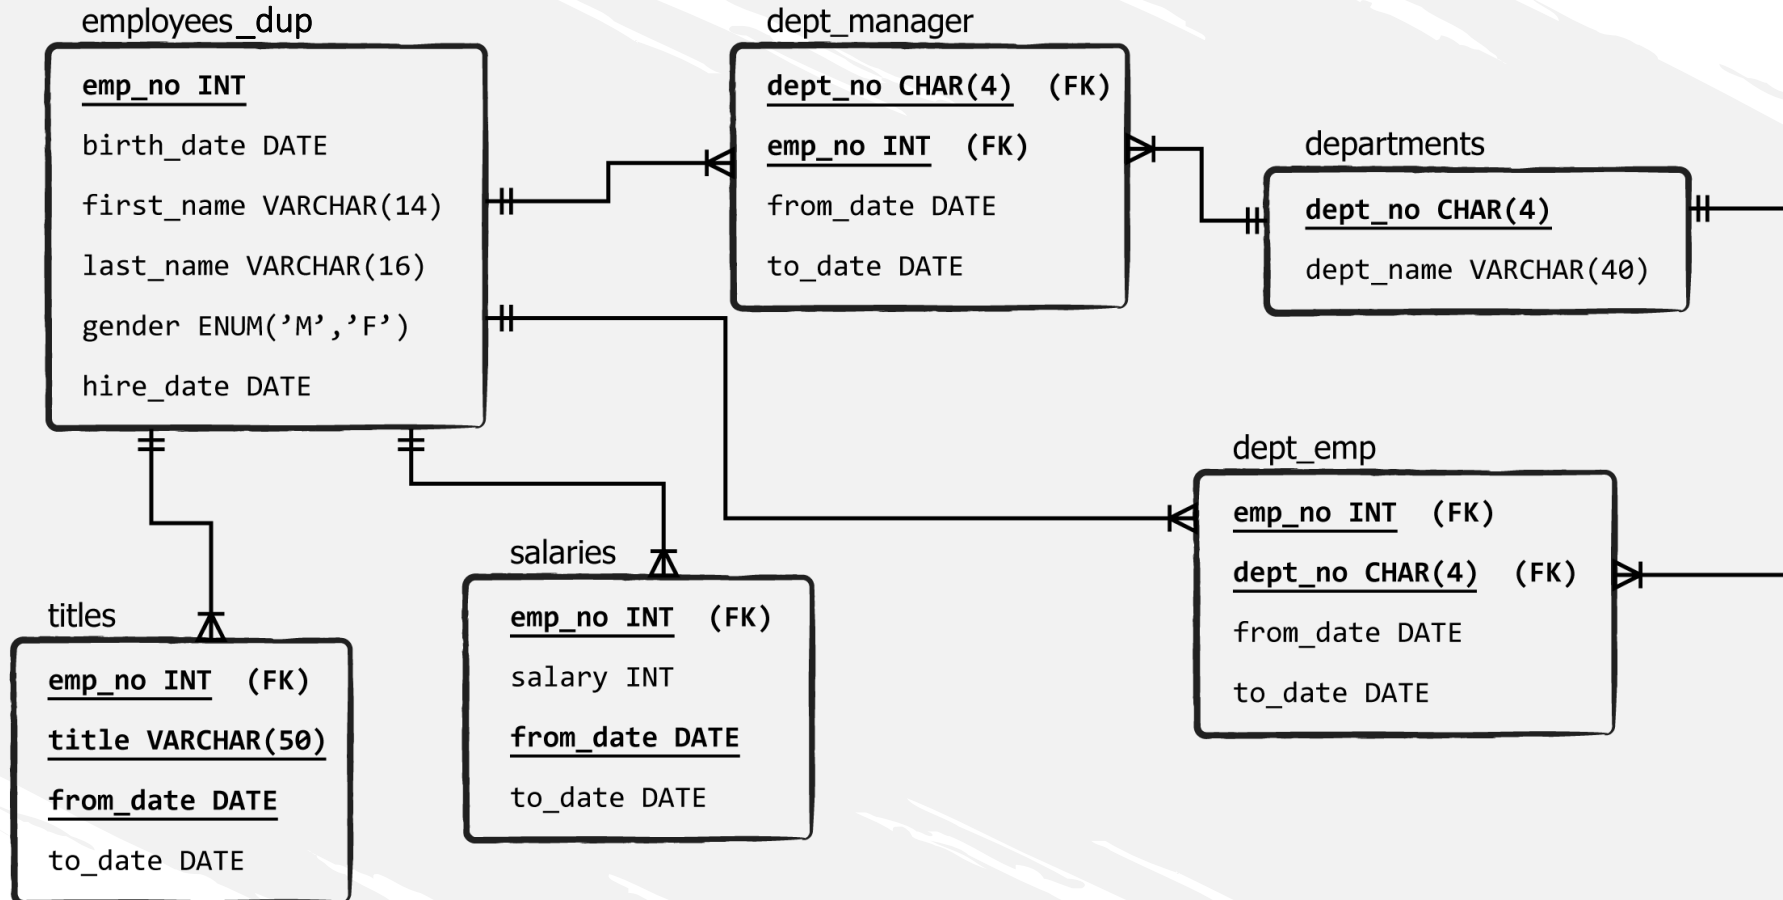
We have to select the same number of columns from each table.

These columns should have the same name, should be in the same order, and should contain related data types.

UNION vs UNION ALL



UNION vs UNION ALL



UNION vs UNION ALL

UNION



SQL

```
SELECT
    N columns
FROM
    table_1
UNION SELECT
    N columns
FROM
    table_2;
```


UNION vs UNION ALL

- when uniting two identically organized tables

UNION vs UNION ALL

- when uniting two identically organized tables
 - UNION displays only distinct values in the output

UNION vs UNION ALL

- when uniting two identically organized tables
 - UNION displays only distinct values in the output
 - UNION ALL retrieves the duplicates as well

UNION vs UNION ALL

- when uniting two identically organized tables
 - UNION displays only distinct values in the output
 - UNION ALL retrieves the duplicates as well

UNION vs UNION ALL

- when uniting two identically organized tables
 - UNION displays only distinct values in the output
 - UNION uses more MySQL resources
 - UNION ALL retrieves the duplicates as well

UNION vs UNION ALL

- when uniting two identically organized tables
 - UNION displays only distinct values in the output
 - UNION uses more MySQL resources (computational power and storage space)
 - UNION ALL retrieves the duplicates as well

UNION vs UNION ALL

- Looking for better results?

UNION vs UNION ALL

- Looking for better results?
 - use UNION

UNION vs UNION ALL

- Looking for better results?
 - use UNION
- Seeking to optimize performance?

UNION vs UNION ALL

- Looking for better results?
 - use UNION
- Seeking to optimize performance?
 - opt for UNION ALL